MASTER THESIS

Thesis submitted in partial fulfillment of the requirements for the degree of Master of Science in Engineering at the University of Applied Sciences Technikum Wien - Degree Program Mechatronics/Robotics

Virtualisierung eines Echtzeit-Betriebssystems zur Steuerung eines Roboters mit Schwerpunkt auf die Einhaltung der Echtzeit

By: Halil Pamuk, BSc

Student Number: 51842568

Supervisor: Sebastian Rauh, MSc. BEng

Wien, August 12, 2024

Declaration

Wien, August 12, 2024   Signature

Kurzfassung

Erstellung einer Echtzeit-Robotersteuerungsplattform unter Verwendung von Salamander OS, Xenomai, QEMU und PCV-521 in der Yocto-Umgebung. Die Plattform basiert auf Salamander OS und nutzt Xenomai für Echtzeit- Funktionen. Dazu muss im ersten Schritt die Virtualisierungsplattform evaluiert werden. (QEMU, Hyper-V, Virtual Box, etc.) Als weiterer Schritt folgt die Anbindung eines Roboters über eine VARAN-Bus Schnittstelle. Das gesamte System wird in der Yocto-Umgebung erstellt und konfiguriert. Das Hauptziel der Arbeit ist es, herauszufinden, wie die Integration von Echtzeit-Funktionen und effizienten Kommunikationssystemen in eine Robotersteuerungsplattform die Reaktionszeit und Zuverlässigkeit von Roboteranwendungen verbessern kann

Schlagworte: Schlagwort1, Schlagwort2, Schlagwort3, Schlagwort4

Abstract

Sections 5.1 and 5.2 demonstrate the inital real-time latency values gathered for bare metal and virtualization.

Contents

1 Introduction

In today's industrial production and automation, robot systems are well established and of crucial importance. Robots must react to their environment and perform time-critical tasks within strict time constraints. Delays or errors can have catastrophic consequences in some cases. Modern operating systems nowadays are categorized into two types: general-purpose operating systems (GPOS) and real-time operating systems (RTOS). The key difference is whether the system is time-critical or not [1]. In a traditional GPOS, such as Windows or Linux-based distributions, a high-priority task cannot interrupt a kernel operation, which makes them not suitable for real-time requirements, as they cannot guarantee deterministic execution times. However, in an RTOS, a high-priority process can interrupt a lower-priority task, even if it is in the middle of a kernel operation. RTOS are specifically designed to react to events within fixed time limits and prioritise the execution of high-priority processes. Many Linux distributions can function as both GPOS and RTOS with kernel modifications.

1.1 Real-Time Operating Systems

The RTOS structure is visualized in Figure 1.The core component of an RTOS that enables real-time capabilities is the kernel. It is responsible for managing system resources, scheduling tasks, and ensuring deterministic behavior. It employs preemptive scheduling mechanisms to allow high-priority tasks to preempt lower-priority tasks, ensuring that time-critical tasks are not delayed. Additionally, RTOS kernels are designed to allocate and manage memory resources efficiently and minimize interrupt latency, which is crucial for real-time applications that require immediate response to external events [2].
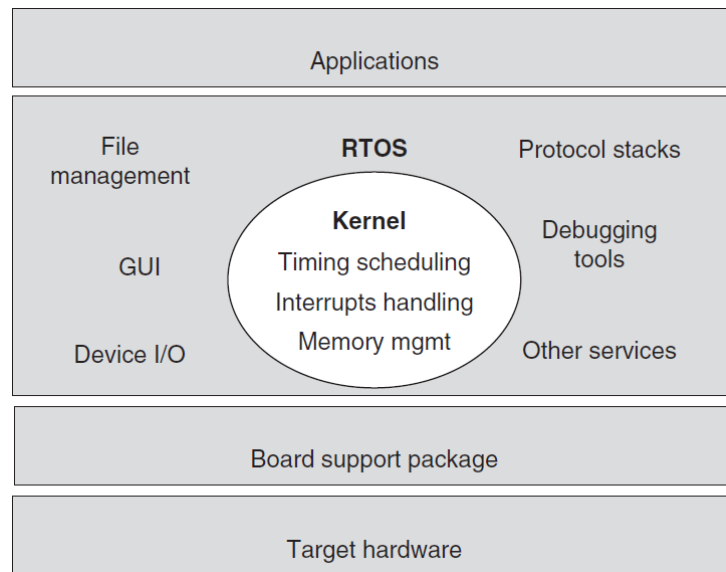
Figure 1: RTOS Structure [2]

In these real-time operating system, task scheduling is based on so-called priority-based preemptive scheduling [3]. Each task in a software application is assigned a priority. A higher priority means that a faster response is required. Preemptive task scheduling ensures a very fast response. Preemptive means that the scheduler can stop a currently running task at any point if it recognizes that another task needs to be executed immediately. The basic rule on which priority-based preemptive scheduling is based is that the task with the highest priority that is ready to run is always the task that must be executed. So if both a task with a lower priority and a task with a higher priority are ready to run, the scheduler ensures that the task with the higher priority runs first. The lower priority task is only executed once the higher priority task has been processed.

Real-time systems are usually classified as either soft or hard real-time systems [4]. The difference lies exclusively in the consequences of a violation of the time limits. Hard real-time is when the system stops operating if a deadline is missed, which can have catastrophic consequences. Soft real-time exists when a system continues to function even if it cannot perform the tasks within a specified time. If the system has missed the deadline, this has no critical consequences. The system continues to run, although it does so with undesirably lower output quality.

1.2 Application context and boundaries

This section defines the boundary conditions within which this master's thesis was composed. This work was written at SIGMATEK GmbH & Co KG [5], a company that researches, develops, produces, and distributes automation technology for industrial machinery and plant engineering. SIGMATEK uses its own custom Linux-based operating system, to be run on their self-manufactured CPUs. This operating system will be referred to as Salamander 4 in this work. The details of this OS are explained in chapter 2.4. Salamander 4 was created with Yocto [6], an open source project that helps develop custom Linux-based systems, as discussed in Section 2.1. Salamander 4 is virtualized through Quick emulator (QEMU), an open-source emulator and virtualizer that allows to run different operating systems on a computer, as described in Section 2.2. Salamander 4 employs hard real-time with Xenomai 3 [7], a real-time framework for Linux that enables deterministic and low-latency performance, detailed in Section 2.3. The goal is to virtualize Salamander 4 and approach the performance of bare metal through real-time performance tunings, as shown in chapter 6, focussing on latency requirements set by SIGMATEK.

1.3 State of the art

1.4 Problem and task definition

1.5 Objective

2 Theoretical Foundations

2.1 Yocto

2.2 QEMU

2.3 Xenomai

Xenomai 3 [7] is a real-time framework that offers two paths to real-time performance. The first approach supplements the Linux kernel with a compact real-time core dubbed Cobalt,

demonstrated in Figure 2. Cobalt runs side-by-side with Linux, but it handles all time-critical activities like interrupt processing and real-time thread scheduling with higher priority than the regular kernel activities.
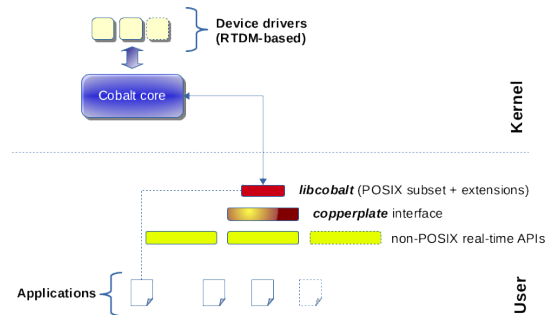
Figure 2: Xenomai Cobalt interfaces [7]

The second approach, called Mercury and shown in Figure 3, relies on the real-time capabilities already present in the native Linux kernel. Often, applications require the PREEMPT-RT extension to augment the mainline kernel's real-time responsiveness and minimize jitter, but this isn't mandatory and depends on the application's specific requirements for responsiveness and tolerance for occasional deadline misses.

Figure 3: Xenomai Mercury interfaces [7]

Salamander 4 uses the Cobalt real-time core with the Dovetail extension, which allows the kernel to handle real-time tasks with low latency.

The main objective of this work is to create a real-time robot control platform that integrates Salamander OS, Xenomai, QEMU and PCV-521 in the Yocto environment.

2.4 Salamander 4

This chapter briefly describes the Salamander 4 operating system by SIGMATEK.

2.4.1 Structure

Salamander 4 is the proprietary operating system of SIGMATEK. It is based on Linux version 5.15.94 and integrates Xenomai 3.2, a real-time development environment [7]. Salamander 4 is a 64-bit system, which refers to the x86_64 architecture. The real-time behaviour is achieved through the use of Symmetric Multi-Processing (SMP) and Preemptive Scheduling (PREEMPT). In addition, it uses IRQPIPE to process interrupts in a way that meets the real-time requirements of the system. The output of the command `uname -a` can be observed in code 1.

```
1 root@sigmatek-core2:~# uname -a
2 Linux sigmatek-core2 5.15.94 #1 SMP PREEMPT IRQPIPE Tue Feb 14 18:18:05 UTC
     2023 x86_64 GNU/Linux
```

Code 1: System information

Salamander 4 is powered by SIGMATEK's CP 841 [8] and is comprised of the following software modules:

- **Operating system**: The operating system in a LASAL CPU manages the hardware and software resources of the system. It is provided in a completely PC-compatible manner, working with a standard PC BIOS.

- **Loader**: The loader is a part of the operating system that is responsible for loading programs from executables into memory, preparing them for execution and then executing them.

- **Hardware classes**: Hardware classes in LASAL represent the different types of hardware components that can be controlled by the LASAL CPU. They provide a way to organize and manage the hardware components in a modular and reusable manner. The graphical hardware editor in LASAL allows for a true-to-detail simulation of the actual hardware.

- **Application**: Applications are developed using LASAL CLASS 2 [9], a solution for automation tasks that supports object-oriented programming and design in compliance with IEC 61131-3.

These modules and the interfaces (indicated by an arrow) between them are shown in Figure 4.

Figure 4: Structure of Salamander 4 CPU

2.4.2 Memory Management

For the sake of completeness, Figure 5 displays the memory management of Salamander 4. LRT stands for Lasal Runtime and creates an execution environment where applications developed using the LASAL Class 2 can run, providing defined real-time functions, data types, and other constructs tailored for real-time programming.
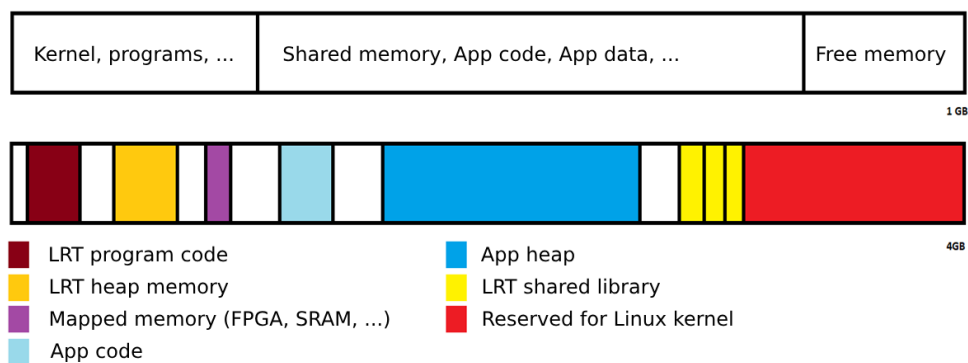
Figure 5: Memory Management

2.5 VARAN-Bus

3 Related Work

4 Methodology

This section describes in detail all the theoretical concepts and boundary conditions as well as practical methods that contributed to achieving the objectives of this master's thesis.

Trace-cmd can be used for tracing the Linux kernel [10]. It can record various kernel events such as interrupts, scheduler decisions, file system activity, function calls in real time. Trace-cmd helped in getting detailed insights into system behaviour and identify reasons for latency.

The data that was recorded by trace-cmd was then fed into Kernelshark, which is a graphical front-end tool [11]. It visualizes the recorded kernel trace data in a readable way on an interactive timeline, which facilitated the process of identifying patterns and correlations between events. By further filtering the displayed events according to specific criteria such as processes, event types or time ranges, the latency issues were analyzed.

Real-time operating system capabilities were provided by Xenomai, which is real-time development framework that extends the Linux kernel [7]. It enables low-latency and deterministic execution of time-critical tasks. Xenomai 3 introduces a dual-kernel approach with a real-time kernel coexisting alongside Linux. A key utility within the Xenomai suite is the latency tool, which benchmarks the timer latency - the time it takes for the kernel to respond to timer interrupts or task activations. The tool creates real-time tasks or interrupt handlers and measures the latency between expected and actual execution times.

The system configuration is shown in Table 1

Table 1: System configuration

CPU	13th Gen Intel(R) Core(TM) i7-13800H
Memory	$2 \times$ 16GB SO-DIMM DDR5-5600 MT/s, 32GB
GPU	NVIDIA RTX A500 Laptop GPU
BIOS	Dell Version 1.12.0
OS	Ubuntu 22.04.4 LTS

Figure 6 is the output of the `lstopo` command and visualizes the hardware nodes of the system, including CPU cores, caches, memory, and I/O devices.
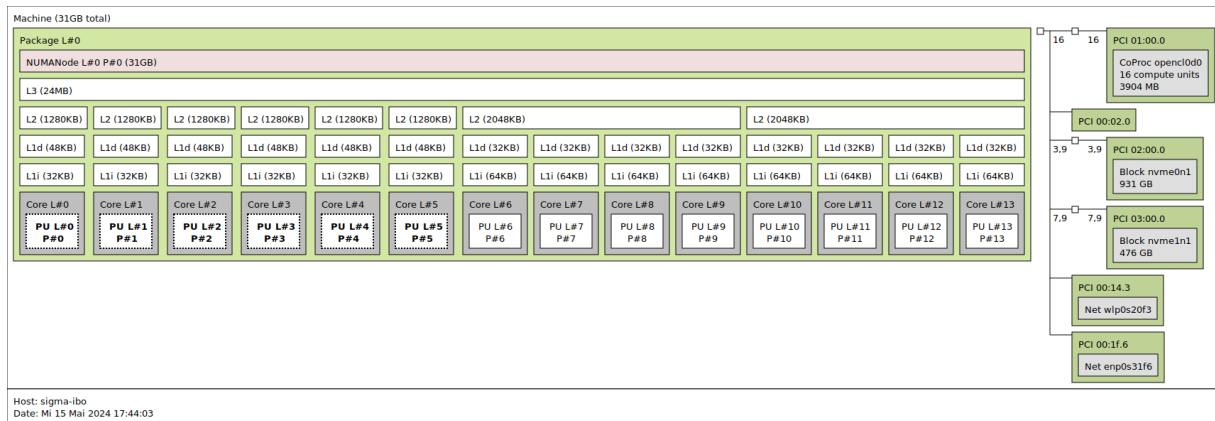
Figure 6: Hardware topology

5 Initial Real-Time Latency

As a starting point, initial latency values of both the bare metal and virtualization versions were measured with the `latency` tool of the xenomai tool suite. Salamander 4 bare metal refers to the proprietary hardware of SIGMATEK used to employ the custom operating system. Salamander 4 virtualization refers to a virtual version of the Salamander 4 hardware platform, achieved through QEMU/KVM. Sections 5.1 and 5.2 specify the details of the measurements for both versions. In the further course, the aim was to bring the latency values of the virtualization closer to those of the hardware and guarantee deterministic and reliable behavior.

5.1 Salamander 4 Bare Metal

The output of the command `uname -a` for Salamander 4 bare metal is shown in Code 2.

```
1 root@sigmatek-core2:~# uname -a
2 Linux sigmatek-core2 5.15.94 #1 SMP PREEMPT IRQPIPE Tue Feb 14 18:18:05 UTC
    2023 x86_64 GNU/Linux
```

Code 2: Salamander 4 bare metal system information

As a reference point, the latency program was executed on Salamander 4 bare metal for a duration of 10 minutes. The complete command used was `latency -h -g gnuplot.txt -T 600`, which runs the latency measurement tool for 600 seconds and prints histograms of min, avg, max latencies in a Gnuplot-compatible format to the file `gnuplot.txt`. Figure 7 shows the gathered latency values in microseconds.

Figure 7: Latency of Salamander 4 bare metal

The statistics obtained from this measurement are provided in Table 2. It gives an overview of the average, maximum, minimum latency, and the standard deviation of the latency values in microseconds.

Table 2: Latency statistics of Salamander 4 bare metal in microseconds

Statistic	Value (µs)
Average Latency	4.06
Maximum Latency	10.71
Minimum Latency	2.82
Standard Deviation	0.85

Figure 8 depicts the variation in latency over the course of said time. Since the data varies strongly, a logarithmic scale was used for both axes.

Figure 8: Variation in latency of Salamander 4 bare metal

5.2 Salamander 4 Virtualization

In addition to providing Salamander 4 on its own hardware, SIGMATEK has also developed a virtualised version of this operating system. It was developed using Yocto, an open-source project that allows customised Linux distributions to be created for embedded systems. Upon generating the necessary files, Yocto provides a QEMU folder with the following components shown in Code 3. QEMU is the environment in which the virtualization runs, as it is an open-source tool for hardware virtualization [12].

```
1    sigma_ibo@localhost:~/Desktop/salamander-image$ ls -1
2    bzImage
3    drive-c
4    ovmf.code.qcow2
5    qemu_def.sh
6    salamander-image-sigmatek-core2.ext4
7    stek-drive-c-image-sigmatek-core2.tar.gz
8    vmlinux
```

Code 3: Contents of QEMU folder for Salamander 4

With the help of the script depicted in Code 4, Salamander 4 is started together with the necessary hardware components in the QEMU environment. This makes it possible to run Salamander 4 on a variety of host systems, regardless of the specific hardware of the host. The following is a description of the components used for the virtualization of Salamander 4.

- **bzImage**: Compressed Linux kernel image that is loaded by QEMU at system start. "bz" stands for big-zipped.

- **drive-c**: Directory serving as C drive for QEMU system, created and filled by qemu_def.sh script.

- **ovmf.code.qcow2**: Firmware file for QEMU that enables UEFI boot process. `OVMF` stands for Open Virtual Machine Firmware, `qcow2` is a format for disk image files used by QEMU, it stands for "QEMU Copy On Write version 2".

- **qemu_def.sh**: Shell script that starts QEMU with required parameters to boot Salamamder 4 OS. It is described in Code 4.

- **salamander-image-sigmatek-core2.ext4**: Disk image of the Salamander 4 OS for the Sigmatek Core 2 platform. It uses the `ext4` file system and serves as the root file system in the QEMU virtual machine, acting as the virtual hard drive.

- **stek-drive-c-image-sigmatek-core2.tar.gz**: Compressed tarball containing a pre-configured environment for the Salamander 4 OS. It is unpacked and sets up the `drive-c/` directory with system and log files in the `qemu_def.sh` script.

- **vmlinux**: Uncompressed Linux kernel image, typically used for debugging.

The inital QEMU script after the custom Yocto build and the starting point for this work is shown in Code 4. This script is used to start QEMU with required parameters to boot Salamamder 4 OS. It will be adjusted in chapter 6 in order to accompany real-time performance tunings.

```
1   #!/bin/sh
2
3   if  [ ! -d drive-c/ ]; then
4        echo "Filling drive-c/"
5        mkdir drive-c/
6        tar -C drive-c/ -xf stek-drive-c-image-sigmatek-core2.tar.gz
7   fi
8
9   exec qemu-system-x86_64 -M pc,accel=kvm -kernel ./bzImage \
10  -m 2048 -drive
        file=salamander-image-sigmatek-core2.ext4,format=raw,media=disk \
11  -append "console=ttyS0 console=tty1 root=/dev/sda rw panic=1
        sigmatek_lrt.QEMU=1 ip=dhcp rootfstype=ext4" \
12  -net nic,model=e1000,netdev=e1000 -netdev bridge,id=e1000,br=nm-bridge \
13  -fsdev local,security_model=none,id=fsdev0,path=drive-c -device
        virtio-9p-pci,id=fs0,fsdev=fsdev0,mount_tag=/mnt/drive-C \
14  -device vhost-vsock-pci,guest-cid=3,id=vsock0 \
15  -drive if=pflash,format=qcow2,file=ovmf.code.qcow2 \
16  -no-reboot -nographic
```

Code 4: QEMU script for starting Salamander 4 virtualization

This script is run on a generic Ubuntu 22.04.4 system, as mentioned previously in chapter 4. The kernel version and other details are presented in Code 5, using the `uname -a` command.

```
1 root@sigmatek-core2:~# uname -a
2 Linux sigma-ibo 6.5.0-35-generic #35~22.04.1-Ubuntu SMP PREEMPT_DYNAMIC Tue
    May  7 09:00:52 UTC 2 x86_64 x86_64 x86_64 GNU/Linux
```

Code 5: Ubuntu 22.04.4 system information

Measuring the latency of the Salamander 4 virtualization with the default QEMU script in Code 4 and no further adjustments for 10 minutes, the following latency values in Figure 9 were collected.

Figure 9: Latency with default settings

The statistics obtained from this measurement are provided in Table 12. It gives an overview of the average, maximum, minimum latency, and the standard deviation of the latency values in microseconds.

Table 3: Latency statistics of default Salamander 4 virtualization in microseconds

Statistic	Value (µs)
Average Latency	46.22
Maximum Latency	707.62
Minimum Latency	15.59
Standard Deviation	52.13

Figure 10 depicts the variation in latency over the course of said time with default settings.

Figure 10: Variation in latency with default settings

Comparing these values to those of bare metal in Figure 11, it is evident that there is a significant initial gap in the statistics. A maximum latency of 707.62 µs is not tolerable for the system and needs to be tuned.

Figure 11: Comparison of variation in latency between hardware and virtualization

6 Real-Time Performance Tuning

In this chapter, the significant initial gap in latency statistics between the virtualized system and the bare metal system is tackled. For this reason, an extensive tuning process is carried out. This involves configurations spanning the BIOS, kernel, host OS, QEMU/KVM virtualization layer, and the Salamander 4 OS itself. The individual configurations will be discussed in detail and the modifications will be justified with a clear explanation. The goal is to bring the latency of the virtualized system closer to that of the bare metal system, thereby ensuring deterministic behavior under real-time constraints.

It is important to mention that a real-time system implies determinism at every component of the system ranging from hardware over the kernel to the software and the application on it. As already established earlier, the guest OS Salamamder 4 runs Xenomai and is hence real-time capable. The host system also needs to be aware of the real-time determinism in order to achieve highest possible reliability. The very step of achieving the goal of reducing latency is to apply the preempt-rt patch [13] to the host system. This patch is a set of modifications to the Linux kernel with the goal of making it fully preemptible. This means, almost all parts of the kernel can be preempted and higher-priority tasks can interrupt lower-priority ones. This significantly reduces the time it takes for high-priority tasks to start executing after an event [14]. The patch also includes support for high-resolution timers and make the system more predictable, which are crucial for real-time applications where the timing of operations must be guaranteed with precise timing and scheduling [15].

Nevertheless, a real-time kernel alone does not make a system truly "real-time" [16]. Additional modifications are required to achieve this. The next sections will deal with these real-time performance tunings. The very first step in this regard is the isolation of a CPU to dedicate it solely to the desired real-time task. CPU isolation and affinity will be covered in detail later in section 6.3.1, but its impact will be briefly demonstrated here.

Figure 12 shows latency of QEMU taskset Salamander4.

Figure 12: Latency taskset

After each tuning, the latency test of Xenomai will be executed to evaluate the improved latency of the system.

After Article, Figure 13

Figure 13: After Article configurations

Table 4: Latency statistics of Salamander 4 virtualization after BIOS configurations in microseconds

Statistic	Value (µs)
Average Latency	17.68
Maximum Latency	32.21
Minimum Latency	8.005
Standard Deviation	6.1

6.1 BIOS Configurations

BIOS stands for Basic Input/Output System. It abstracts the hardware and enables basic functions of a computer during the booting process, such as starting the operating system and loading other software. Since the BIOS is embedded very deep, its configuration can significantly influence the real-time performance of the system. Table 5 illustrates the specific BIOS settings that have been adjusted for the purpose of real-time performance.

Table 5: BIOS Configurations

Option	Status
Hyper Threading	Disabled
Intel SpeedStep®	Disabled
Intel® Speed Shift Technology	Disabled
C States	Disabled
VT-d	Enabled

In the following, these settings along with their impact on system latency are briefly described.

- **Hyper Threading**: Hyper-Threading allows CPUs to process two threads simultaneously instead of just one. When it is enabled on the host, this allows the parallelisation of tasks and increases performance. However, in a real-time system like the guest Salamander 4, this can lead to increased latencies due to contention between threads. In order to ensure more deterministic behavior in the guest, it is disabled on the host.

- **Intel SpeedStep**: This technology dynamically adjusts the clock speed of the CPU based on workload. These dynamic changes can lead to unpredictable latencies in a real-time system. It is also disabled on the host to maintain a constant CPU speed.

- **Intel® Speed Shift Technology**: Similar to SpeedStep, Speed Shift allows the processor to directly control its frequency and voltage. This can lead to unpredictable latencies, too. Hence, it is also disabled on the host.

- **C States**: These are low-power idle states where the clock frequency and voltage of the CPU are reduced. Transitioning between C-states can cause variable latencies. To prevent this from happening, C-states are disabled on the host.

- **VT-d**: Direct access to physical devices from within virtual machines is possible when VT-d is enabled on the host. This can help reduce latencies associated with I/O operations in the virtual machine. It is therefore enabled on the host.

6.2 Kernel Configurations

The kernel command-line parameters are shown in Code 6 below.

```
1      GRUB_CMDLINE_LINUX="isolcpus=4 rcu_nocbs=4 rcu_nocb_poll nohz_full=4
           nohz=on default_hugepagesz=1G hugepagesz=1G hugepages=8
           intel_iommu=on rdt=l3cat nmi_watchdog=0 idle=poll clocksource=tsc
           tsc=reliable audit=0 skew_tick=1 intel_pstate=disable
           intel.max_cstate=0 intel_idle.max_cstate=0 processor.max_cstate=0
           processor_idle.max_cstate=0 nosoftlockup no_timer_check
           nospectre_v2 spectre_v2_user=off kvm.kvmclock_periodic_sync=N
           kvm_intel.ple_gap=0 irqaffinity=0"
```

Code 6: Kernel Configuration

In the following, these settings along with their impact on system latency are briefly described.

- **isolcpus=4**: Isolates CPU 4 from the general scheduler, meaning no process will be scheduled to run on this CPU unless it is explicitly assigned. CPU isolation is explained in detail in section 6.3.1.

- **rcu_nocbs=4**: The Linux kernel uses a synchronization mechanism called RCU, or Read-Copy-Update. It lets writers update the data in a way that guarantees readers will always see the same version while enabling multiple readers to access shared data without locks. The RCU subsystem uses callback functions that need to be invoked once readers are done with the data they accessed. By default, these callbacks are handled by the CPUs that executed the read-side critical sections. This parameter offloads RCU callback handling from CPU 4 to other CPUs. CPU 4 remains dedicated to high-priority tasks which helps in reducing latency.

- **rcu_nocb_poll**: This is used together with `rcu_nocbs` and causes the system to actively poll for RCU callbacks to invoke, instead of waiting for the next RCU grace period. This reduces latency.

- **nohz_full=4**: Makes CPU core 4 "tickless", meaning the kernel tries to avoid sending periodic scheduling-clock interrupts to the CPU when there are no runnable tasks. This lowers latency by reducing unnecessary wake-ups but may increase power consumption because the CPU is not able to enter a low-power state when idle. Additionally, timer interrupts cannot be fully eliminated because certain events, such as incoming interrupts or task activations, can still cause the kernel to send timer interrupts to the tickless CPU.

- **nohz=on**: Sets all CPUs to tickless mode system-wide.

- **default_hugepagesz=1G, hugepagesz=1G, hugepages=8**: Huge pages are large contiguous areas of memory that can be used by applications and the kernel, instead of the traditional 4KB small pages. The default huge page size is set to 1GB, and 8 huge pages of 1GB size are reserved at boot. This pre-allocation makes sure that these large memory regions are available to be used by the kernel or applications, without having to dynamically allocate and potentially fail.

- **intel_iommu=on**: Enables Intel's IOMMU (Input/Output Memory Management Unit), which connects a DMA (Direct Memory Access)-capable I/O bus, such as graphics cards and network adapters, to the main memory. It can enhance device performance by allowing these devices to directly access and use memory, which is especially helpful when these devices are virtualized.

- **rdt=l3cat**: Activates the L3 CAT (L3 Cache Allocation Technology) feature of Intel's RDT (Resource Director Technology). Unlike L1 and L2 caches, where each core has its fixed capacity, L3 cache is a shared pool among multiple cores. L3 CAT is a mechanism that controls the amount of L3 cache that a process can use. By controlling cache allocation, it can prevent a single process from monopolizing the L3 cache, which is particularly beneficial in virtualized environments, where multiple virtual machines share the same physical host.

- **nmi_watchdog=0**: Disables the NMI (Non-Maskable Interrupt) watchdog, which is a debugging feature of the Linux kernel. It works by periodically generating non-maskable interrupts. If the system does not respond to these interrupts within a certain timeframe, the NMI watchdog concludes that the system has hung and generates a system dump for debugging. This constant monitoring consumes CPU cycles and can introduce undesirable latency in real-time systems.

- **idle=poll**: Changes the CPU's idle loop behavior to active polling. Instead of entering a low-power state when idle, the CPU continuously polls for new tasks. This can reduce task start latency in real-time systems, but it increases power consumption.

- **clocksource=tsc, tsc=reliable**: TSC (Time Stamp Counter) is a high-resolution timer provided by most x86 processors that counts the number of CPU cycles since it was last reset. Accurate timekeeping is crucial, particularly for real-time systems. These parameters set the clocksource to TSC and mark it as a reliable source of timekeeping, meaning it increments at a consistent rate and does not stop when the processor is idle.

- **audit=0**: Disables the Linux audit system. When it is enabled, it generates log entries for security-relevant events, which is a slow operation since they are written to disk. If there are a large number of such events, the audit system can consume significant CPU time and I/O bandwidth which could lead to higher latency.

- **skew_tick=1**: Enables a mode in the Linux kernel that reduces timer interrupt overhead. Normally, timer interrupts happen simultaneously on all CPUs, resulting in all CPUs to exit their low-power states at once. This can lead to increased contention for system resources. When enabled, the kernel offsets the timer interrupts on different CPUs, spreading them out over time.

- **intel_pstate=disable**: Disables the Intel P-state driver, which is a part of the Linux kernel that handles power management for Intel CPUs. It controls the frequency of the CPU by

scaling it up when demand is high and scaling it down to save power when demand is low. This dynamic frequency scaling is disabled because it leads to increased latencies for real-time systems.

- **intel.max_cstate=0, intel_idle.max_cstate=0, processor.max_cstate=0, processor_idle.max_cstate=0**: These parameters disable deeper C-states (CPU power saving states). Normally, when a CPU is idle, it can enter various C-states, with higher-numbered states representing deeper sleep states that save more power but take longer to wake up from. In real-time systems, these wake-up delays can be problematic. Disabling them leads keeps the CPUs ready to respond quickly to new tasks and helps in reducing latency.

- **nosoftlockup**: Disables the soft lockup detector in the Linux kernel. A soft lockup is when a CPU is busy executing kernel code for a long period of time without giving other tasks a chance to run. Especially threads with SCHED_FIFO policy occupy the CPU for an extensive duration. This is detected and reported by the soft lockup detector, hence it is disabled to prevent these unnecessary warnings.

- **no_timer_check**: Disables the check for broken timer interrupt sources. Broken timer interrupt sources are problems with hardware or software that prevent timer interrupts from working as intended. Such a timer may not generate interrupts at the expected rate or at all. The kernel skips the checks for these broken timer interrupt sources, which can cause unnecessary overhead in a real-time system where every CPU cycle counts.

- **nospectre_v2, spectre_v2_user=off**: These parameters disable mitigations for the Spectre v2 vulnerability. Spectre v2 is a hardware vulnerability that affects many modern microprocessors and can allow malicious programs to access sensitive data they are not supposed to. While this is necessary for security, it has an impact on performance and should be turned off in controlled environments where the risk of exploitation is low.

- **kvm.kvmclock_periodic_sync=N, kvm_intel.ple_gap=0**: These are KVM (Kernel-based Virtual Machine) related parameters. They disable the periodic synchronization of the kvmclock and set the gap between PLE (Pause Loop Exiting) events to 0. The kvmclock is a paravirtualized clock source provided by KVM to its guest OS and disabling this reduces latency introduced by clock synchronization. By setting the gap to 0, the virtual machine exits the pause loop immediately, which can reduce latency in spinlock-intensive workloads.

- **irqaffinity=0**: Sets the default Interrupt Request affinity to none. This means that no CPU core is preferred over another for handling IRQs. Instead, it lets the operating system decide how to distribute these IRQs across all the CPUs. Section 6.3.3 dives deeper into Interrupt Request affinity.

After the BIOS and Kernel configurations, Figure 14

Figure 14: After BIOS and Kernel configurations

Table 6: Latency statistics of Salamander 4 virtualization after Kernel configurations in microseconds

Statistic	Value (µs)
Average Latency	14.00
Maximum Latency	21.694
Minimum Latency	6.351
Standard Deviation	1.44

6.3 Host OS Configurations

The host OS needs to provide an environment where the guest OS can operate in real-time. This entails a number of host-side adjustments to reduce interruptions and reduce latency. In the following, a detailed overview of these configurations and their impact on the real-time performance of the guest OS is provided.

6.3.1 CPU affinity and isolation

Isolating CPUs involves removing all user-space threads and unbound kernel threads since bound kernel threads are tied to specific CPUs and hence cannot be moved. For CPU isolation, the `isolcpus` function was used to isolate a performance CPU from the general scheduling algorithms of the operating system. This means that the isolated CPUs will not be used for regular task scheduling, allowing them to be dedicated for the real-time specific tasks. However, the isolcpus function only isolates at the user level and does not affect kernel tasks. Consequently, these kernel tasks and interrupts can still utilize the CPU [17], including systemd services. To prevent systemd services from running on an isolated CPU, the CPU affinity can be set the with line `CPUAffinity=0 1 2 3 5 6 7 8 9 10 11 12 13` in the `/etc/systemd/system.conf` file to indicate the CPUs that systemd-services are allowed to run on. Every CPU other than the isolated CPU 4 is allowed. Code 7 shows the user and kernel tasks that run on CPU 4. After the isolation, user tasks other than the QEMU process have been removed from running on this CPU. Only few per-CPU kernel threads that are tied to this CPU still take CPU time.

```
1    sigma_ibo@sigma-ibo:~$ cat /sys/devices/system/cpu/isolated
2    4
3    sigma_ibo@sigma-ibo:~$ ps axHo psr,pid,lwp,args,policy,nice,rtprio |
        awk '$1 == 4'
4    4      38       38 [cpuhp/4]                  TS     0       -
5    4      39       39 [idle_inject/4]            FF     -      50
6    4      40       40 [migration/4]              FF     -      99
7    4      41       41 [ksoftirqd/4]              TS     0       -
8    4      42       42 [kworker/4:0-events]       TS     0       -
9    4      43       43 [kworker/4:0H-kblockd]     TS   -20       -
10   4     153      153 [kworker/4:1-events]       TS     0       -
11   4   81649    81649 qemu-system-x86_64 -M pc,ac TS    0       -
12   4   81649    81654 qemu-system-x86_64 -M pc,ac TS    0       -
13   4   81649    81676 qemu-system-x86_64 -M pc,ac TS    0       -
14   4   81649    81702 qemu-system-x86_64 -M pc,ac TS    0       -
15   4   81649    82185 qemu-system-x86_64 -M pc,ac TS    0       -
16   4   81649    82187 qemu-system-x86_64 -M pc,ac TS    0       -
17   4   82134    82134 [kworker/4:1H-kblockd]     TS   -20       -
```

Code 7: User and Kernel Tasks

6.3.2 KVM entry and KVM exit

Upon isolating a CPU to the QEMU process, it was anticipated that the guest would utilize nearly 100% of the CPU's capacity, with minimal to no intervention from the host. However, the isolcpus function only isolates CPUs from the scheduler, meaning that user-space processes won't be scheduled on these CPUs, but kernel tasks and interrupts can still utilize the CPU.This led to the investigation of the causes for the observed high and inconsistent latency. The guest

operates within the `kvm_entry` and `kvm_exit` events of the host. Kernelshark revealed a high frequency of `kvm_exit` events, indicating that the guest frequently relinquishes control of the CPU back to the host. This frequent switching hinders the guest's ability to run continuously, thereby increasing the virtualization latency.

To further understand this, trace-cmd was employed to trace various events in the host-guest communication, including the reasons for these events. Specifically, the causes for `kvm_exit` events were analyzed. The command sudo trace-cmd record -e all -A @3:823 –name Salamander4 -e all was executed on the host for a duration of 5 seconds. The results in Figure **??** were obtained. Additionally, Table 7 provides a short description of the observed `kvm_exit` events.

Exit Reason	Description
APIC_WRITE	Triggered when the guest writes to its APIC.
EXTERNAL_INTERRUPT	Triggered by external hardware interrupts.
HLT	Triggered when the guest executes the HLT instruction.
EPT_MISCONFIG	Triggered by a misconfiguration in the EPT.
PREEMPTION_TIMER	Triggered when the host's preemption timer expires.
PAUSE_INSTRUCTION	Triggered when the PAUSE instruction is executed.
EPT_VIOLATION	Triggered by a violation of the EPT permission settings.
IO_INSTRUCTION	Triggered when the guest executes an I/O instruction.
EOI_INDUCED	Triggered when an EOI signal is sent to the APIC.
MSR_READ	Triggered when the guest reads from a MSR.
CPUID	Triggered when the guest executes the CPUID instruction.

Table 7: Description of kvm_exit reasons

In the process of analyzing the `kvm_exit` events, several reasons for these exits were identified. The most frequent among these were the `APIC_WRITE` and `HLT` events. The former is initiated when the guest writes to its Advanced Programmable Interrupt Controller (APIC), a component of the CPU that manages hardware interrupts. The latter occurs when the guest executes the HLT instruction, effectively halting the CPU until the next external interrupt is fired. Other significant but less frequent events included `EXTERNAL_INTERRUPT` and `IO_INSTRUCTION`. These events are indicative of the guest's interaction with hardware devices and its execution of I/O operations. Events such as `EPT_MISCONFIG` and `PREEMPTION_TIMER` were also noted. These could potentially signal issues with memory management and the host's scheduling of the guest. While events like `PAUSE_INSTRUCTION`, `EPT_VIOLATION`, `EOI_INDUCED`, `MSR_READ`, and `CPUID` were the least frequent, they still provide valuable insights into the guest's behavior and the host-guest interaction.

When the script is started from the host, the QEMU process can be scheduled to run on any available core, as it is not bound to a specific CPU core. This means that the QEMU process may frequently switch between different cores, leading to an increase in latency. As the goal was to reduce latency in the guest, the first step was to isolate a CPU of the host and dedicate it solely to the QEMU process, so that it cannot be used for other tasks on the user level. Without CPU isolation, context switches take place at the operating system level and not at the hypervisor level. This explains why there are fewer `kvm_exit` events. However, this will also, as previously shown, lead to higher latency, as context switches at the operating system level generally take longer than a `kvm_exit` and `kvm_entry`. When the CPU was dedicated to the QEMU process, on the other hand, there was a significant increase in `kvm_exit` events. This is because every context switch takes place at the hypervisor level. Nevertheless, lower latency was achieved thereby, as the qemu process is no longer influenced by the CPU scheduling of the operating system.

6.3.3 Interrupt Affinity

Once the CPUs were isolated, interrupt requests handling was the next step. The purpose of interrupt requests is to inform the CPU to stop working on a certain job and start working on another. This allows hardware devices to communicate with the CPU through frequent context switches, which can introduce latency in real-time systems. The `/proc/interrupts` file can be monitored using the command `watch -d -n 1 cat /proc/interrupts` to observe changes in the interrupt requests handled by each CPU in real-time. The IRQs needed to be removed from the isolated CPU by manipulating the `/proc/irq/<IRQ>/smp_affinity` files. The value in the `smp_affinity` file is a bitmask in hexadecimal format where each bit corresponds to a CPU. The least significant bit (LSB) on the right corresponds to the first CPU (CPU0), and the most significant bit (MSB) on the left corresponds to the last CPU (CPU13). When there are 14 CPUs available, the default value for `smp_affinity` would be 3FFF. Removing CPU 4 out of this bitmask would be setting bit five to zero, resulting in 3FEF (11 1111 1110 1111). The python script in Code 8 was written to show a table of the distribution of interrupt requests across each CPU. By changing the values of the `smp_affinity` files, the assignment of IRQs to CPUs is controlled.

```
1    import os
2    import pandas as pd
3    from tabulate import tabulate
4
5    # Get the number of CPUs
6    num_cpus = os.cpu_count()
7
8    # Initialize a dictionary to store the CPUs for each IRQ
9    irqs = {}
10
11   # Iterate over each IRQ
12   for irq in os.listdir('/proc/irq'):
13       # Check if the smp_affinity file exists for this IRQ
14       if os.path.isfile(f'/proc/irq/{irq}/smp_affinity'):
15           # Read the current smp_affinity
16           with open(f'/proc/irq/{irq}/smp_affinity', 'r') as f:
17               affinity = int(f.read().strip(), 16)
18           # Initialize an empty list to store the CPUs for this IRQ
19           cpus = []
20           # Iterate over each CPU
21           for cpu in range(num_cpus):
22               # Check if the bit for the current CPU is set
23               if ((affinity & (1 << cpu)) != 0):
24                   # Add the CPU to the list for this IRQ
25                   cpus.append(cpu)
26           # Sort the list of CPUs
27           cpus.sort()
28           # Add the list of CPUs to the dictionary for this IRQ
29           irqs[irq] = cpus
30
31   # Create a DataFrame to store the table
32   df = pd.DataFrame(index=sorted(irqs.keys(), key=int),
        columns=range(num_cpus))
33
34   # Fill the DataFrame with 'x' where a CPU is assigned to an IRQ
35   for irq, cpus in irqs.items():
36       for cpu in cpus:
37           df.loc[irq, cpu] = 'x'
38
39   # Replace NaN values with empty strings
40   df.fillna('', inplace=True)
41
42   # Print the table in pipe format
43   print(tabulate(df, headers='keys', tablefmt='pipe', showindex=True))
44
45   # Convert the DataFrame to a markdown table
46   markdown_table = df.to_markdown()
47
48   # Write the markdown table to a file
49   with open('table_CPU_IRQ.md', 'w') as f:
50       f.write(markdown_table)
```

Code 8: Check distribution of interrupt requests across each CPU

However, as the proc filesystem resets to its default state after each reboot, manually changing numerous IRQ files is tedious and time-consuming. This process was automated using the shell script in Code 9, which is executed after every reboot. Additionally, the CPU affinity of IRQ threads of NVME (Non-Volatile Memory Express, a type of SSD storage) can be set away from CPU 4 to avoid impacting real-time workloads. This can be done through finding out the respective process IDs through `ps -e | grep irq/.*nvme` and then executing `sudo taskset -a -p -c 0 <PID>`.

```bash
1    #!/bin/bash
2
3    # Check if a command-line argument is provided
4    if [ -z "$1" ]; then
5        echo "Please provide a CPU number as a command-line argument."
6        exit 1
7    fi
8
9    # Get the CPU number from the command-line argument
10   CPU=$1
11
12   # Define the mask values
13   declare -A mask_values
14   mask_values=( [0]="3ffe" [1]="3ffd" [2]="3ffb" [3]="3ff7" [4]="3fef"
         [5]="3fdf" [6]="3fbf" [7]="3f7f" [8]="3eff" [9]="3dff" [10]="3bff"
         [11]="37ff" [12]="2fff" [13]="17ff")
15
16   # Run the check_smp_affinity.sh script and get the IRQs
17   IRQs=$(./check_smp_affinity.sh $CPU | grep -o '[0-9]\+')
18
19   # Initialize an empty array to store the IRQs that could not be removed
20   failed_IRQs=()
21   # Initialize an empty array to store the IRQs that were successfully
         removed
22   succeeded_IRQs=()
23
24   # Loop over the IRQs
25   for IRQ in $IRQs; do
26       # Try to change the smp_affinity
27       echo ${mask_values[$CPU]} | sudo tee /proc/irq/$IRQ/smp_affinity >
           /dev/null 2>&1
28       # If the command failed, add the IRQ to the failed_IRQs array
29       if [ $? -ne 0 ]; then
30           failed_IRQs+=($IRQ)
31       else
32           succeeded_IRQs+=($IRQ)
33       fi
34   done
35
36   # Check if there were any failed IRQs
37   if [ ${#failed_IRQs[@]} -ne 0 ]; then
38       echo "IRQs ${failed_IRQs[@]} could not be removed from CPU $CPU."
39   fi
40
41   # Check if there were any successful IRQs
42   if [ ${#succeeded_IRQs[@]} -ne 0 ]; then
43       # Remove the first entry from the succeeded_IRQs array
44       succeeded_IRQs=("${succeeded_IRQs[@]:1}")
45       echo "IRQs ${succeeded_IRQs[@]} were removed from CPU $CPU."
46   fi
```

Code 9: Change IRQ assignment of a CPU

6.3.4 RT-priority

Having a real-time kernel itself is a crucial part of achieving deterministic behavior, but not enough to take full advantage of the real-time capabilities. One key aspect of this are real-time priorities, thoroughly explained by Richard Weinberger in [18]. In essence, Table 8 lists minimum and maximum priorities for different scheduling policies in Linux.

Table 8: Minimum and maximum priorities for different scheduling policies

Scheduling Policy	Min Priority	Max Priority
SCHED_OTHER	0	0
SCHED_FIFO	1	99
SCHED_RR	1	99
SCHED_BATCH	0	0
SCHED_IDLE	0	0
SCHED_DEADLINE	0	0

`SCHED_FIFO` allows deterministic, high-priority execution of critical tasks without being pre-empted by lower-priority processes. The virtual machine can either be started with `chrt -f <PRIO>` or adjusted at a later point with `chrt -f <PRIO> <PID>`. It is also important that there are no other unexpected real time processes running on the system concurrently.

6.3.5 Disable RT throttling

If a real-time task consumes 100% of the CPU time, the system may become unresponsive as a whole. This happens because an RT process is constantly using the CPU and the Linux scheduler will not schedule other non-RT processes in the meantime. To prevent complete system lockups, the kernel has a function to throttle RT processes if they consume 0.95 seconds out of every 1 second of CPU time. It does this by pausing the process for the remaining 0.05 seconds, which is not desired because this could result in missed deadlines. RT throttling can be disabled by writing the value -1 to the `/proc/sys/kernel/sched_rt_runtime_us` file. This change also needs to be made permanent because the proc filesystem resets to its default state after each reboot. For this purpose, the line `kernel.sched_rt_runtime_us = -1` can be appended to the end of the `/etc/sysctl.conf` file, which is read at boot time and used to configure kernel parameters. This reduces the potential for missed deadlines.

6.3.6 Disable timer migration

Timer migration allows timers to be moved from one CPU to another, which means the kernel can balance load across multiple CPUs. In a real-time system, this can introduce latency and jitter. To disable timer migration, the value "0" needs to be written to the `/proc/sys/kernel/timer_migration` file. This change also needs to be made permanent by writing the line `kernel.timer_migration = 0` to the `/etc/sysctl.conf` file. This reduces the amount of context switches and interrupts.

6.3.7 Set Device Driver Work Queue

The device driver work queue allows time-consuming tasks to be offloaded to be processed later in a separate kernel thread. By setting the work queue away from CPU 4 to another CPU, it is free to handle real-time tasks without being interrupted by these work queue tasks. This is done by specifying a bitmask to exclude CPU 4 in the files `/sys/devices/virtual/workqueue/cpumask` and `/sys/bus/workqueue/devices/writeback/cpumask`.

6.3.8 Disable RCU CPU stall warnings

As already mentioned in section 6.2, the Linux kernel uses RCU as a synchronization mechanism for reading from and writing to shared data. An RCU CPU stall in the Linux kernel can occur due to several reasons. These include a CPU looping in an RCU read-side critical section, a CPU looping with interrupts disabled, a CPU looping with preemption disabled, or a CPU not getting around to less urgent tasks, known as "bottom halves". The number of seconds the kernel should wait before checking for stalled CPUs and reporting a stall warning can be set via the `/sys/module/rcupdate/parameters/rcu_cpu_stall_timeout` file. These warnings can be suppressed alltogether to reduce the potential for increased latency by writing "1" to the `/sys/module/rcupdate/parameters/rcu_cpu_stall_suppress` file. This setting is also not persistent across reboots, so the command needs to be added to a startup script.

6.3.9 Stop Certain Services

Services like irqbalance.service, thermald.service, and wpa_supplicant.service, as presented in Table 9 can be further sources for random latency and unnecessary overhead. Stopping these services through `sudo systemctl stop <SERVICE>` means they will not be able to interrupt the CPU with their tasks.

Table 9: Description of Services

Service	Description
irqbalance.service	Distributes hardware interrupts across CPUs
thermald.service	A daemon that prevents overheating
wpa_supplicant.service	A service for wireless network devices

6.3.10 Disable Machine Check

Machine checks report hardware errors and these checks can cause interruptions and increase latency. Hence it is best to disable them in real-time scenarios by writing a "0" to the `/sys/devices/system/machinecheck/machinecheck0/check_interval` file.

6.3.11 Boot into text-based environment

The Graphical User Interface (GUI) consumes great amounts of system resources. It often runs different background processes that are not essential for real-time systems and hence increase latency. These resources can be freed up by switching to a text-based environment. This way, processing power and memory can be allocated to critical real-time tasks which in turn can lead to lower latency and deterministic behavior. The command `systemctl set-default multi-user.target` and a following reboot allow for booting into a text-based environment. For the sake of completeness, the graphical interface can be brought back through the command `systemctl set-default graphical.target` and a following reboot.

After the Host configurations, Figure 15

Figure 15: After Host configurations

Table 10: Latency statistics of Salamander 4 virtualization after Host configurations in microseconds

Statistic	Value (µs)
Average Latency	12.61
Maximum Latency	17.041
Minimum Latency	7.872
Standard Deviation	1.8

6.4 QEMU/KVM Configurations

KVM allows the guest OS to run directly on the hardware, bypassing the need for traditional emulation which can introduce delays. QEMU, when used with KVM, provides hardware-assisted virtualization, which also lowers latency in the guest OS.

6.4.1 Tune lapic timer advance

The Local Advanced Programmable Interrupt Controller (LAPIC) is a built-in timer that handles the delivery of interrupts to the CPU. It generates interrupts at a rate. This rate can be

tuned in the `/sys/module/kvm/parameters/lapic_timer_advance_ns` file to reduce the frequency of interrupts and therefore decrease the latency of the guest VM. The default value is "-1", which means that the kernel will automatically calculate an appropriate advance for the timer. Here, it is set to the value "7500". Hence,the timer interrupt will be delivered 7500 nanoseconds earlier than it is actually due. This gives the VM more time to handle it.

6.4.2 Set QEMU options for real-time VM

QEMU provides several options that can be used to improve the real-time performance of the guest VM. Table 11 briefly explains these options.

Table 11: QEMU options for real-time performance

QEMU Option	Description
`-object memory-backend-ram,` `id=ram0,size=4G,prealloc=on`	Locks the memory of the VM to 4GB and prevents it from being swapped out to disk
`-mem-prealloc` `-mem-path /dev/hugepages/`	Enables the use of hugepages and improves memory access

Code 10 shows the final QEMU script used to start the Salamander 4 virtualization, including these options.

```
1      #!/bin/sh
2
3      if  [ ! -d drive-c/ ]; then
4            echo "Filling drive-c/"
5            mkdir drive-c/
6            tar -C drive-c/ -xf stek-drive-c-image-sigmatek-core2.tar.gz
7      fi
8
9      exec taskset -c 4 qemu-system-x86_64 -M pc,accel=kvm -kernel ./bzImage \
10     -m 2048 -drive
            file=salamander-image-sigmatek-core2.ext4,format=raw,media=disk \
11     -append "console=ttyS0 console=tty1 root=/dev/sda rw panic=1
            sigmatek_lrt.QEMU=1 ip=dhcp rootfstype=ext4 schedstats=enable nohlt
            idle=poll quiet xeno_hal.smi=1 xenomai.smi=1 threadirqs" \
12     -net nic,model=e1000,netdev=e1000 -netdev bridge,id=e1000,br=nm-bridge \
13     -fsdev local,security_model=none,id=fsdev0,path=drive-c -device
            virtio-9p-pci,id=fs0,fsdev=fsdev0,mount_tag=/mnt/drive-C \
14     -device vhost-vsock-pci,guest-cid=3,id=vsock0 \
15     -drive if=pflash,format=qcow2,file=ovmf.code.qcow2 \
16     -object memory-backend-ram,id=ram0,size=4G,prealloc=on \
17     -mem-prealloc -mem-path /dev/hugepages \
18     -no-reboot -nographic
```

Code 10: Tuned QEMU script for starting Salamander 4 virtualization

After the QEMU configurations, Figure 16

Figure 16: After QEMU configurations

Table 12: Latency statistics of Salamander 4 virtualization after QEMU configurations in microseconds

Statistic	Value (µs)
Average Latency	16.26
Maximum Latency	17.134
Minimum Latency	14.532
Standard Deviation	0.3

6.5 Guest OS Configurations

6.6 Other configurations

There are other configurations that may be relevant depending on the case. Linux Kernel Developer Steven Rostedt explains all relevant aspects of a real-time system that must be considered in [19] and gives insight for finding sources of latency on the linux system in [20]. A Checklist for Writing Linux Real-Time Applications is provided by John Ogness in [21]. Every layer of the system stack must be deterministic to ensure predictable and reliable latency, including hardware, operating system, middleware and drivers, and the application software. [22] and [23] describe the process of writing hard real time Linux programs using the real time preemption patch in great detail. Various hardware and software tunings are mentioned in [24] and [25].

7 Real-Time Robotic Application

This chapter compares the latency of the Salamander 4 operating system before and after the real-time performance tunings with the latency of Salamander 4 running on bare metal hardware. This comparison is done with the aid of a program that will be explained shortly. The goal is to understand how closely the performance of the virtualization can match that of the bare metal. The experimental setup includes a six-axis mini-robot, illustrated in Figure 17.

Figure 17: Mini-robot of the experiment

The drive system of the robot arm consists of six MG996R digital servo motors [26], equipped with a metal gearbox. The motor is able to rotate in a range of approximately 180 degrees and its position can be controlled with a high degree of accuracy. Each servo motor has three wires that need to be connected as shown in Figure 18.

Figure 18: MG996R Servo Motor [26]

To drive the motor, it has to be powered using the red and brown wires and can be controlled by sending PWM signals to the orange wire. In this experiment, this PWM signal is generated by the proprietary PW 022 pulse width module of Sigmatek [27]. Its connector layout is illustrated in Figure 19.

Modul Status (green)
User (yellow)
DC OK (green)

PW 022

X1
24 — Supply PWM Outputs
0V — GND
O1 — PWM Output 1
O2 — PWM Output 2

X2
nc — do not use
nc — do not use
0V — GND
0V — GND

SIGMATEK

Figure 19: PW 022 pulse width module connector layout [27]

The module has two +24 V switching PWM outputs with an adjustable frequency for controlling inductive loads. Since the mentioned servo motors operate between 4.8 volts and 7.2 volts [26], this voltage needed to be reduced through resistors before supplying it to a servo motor. For this purpose, two resistors with resistances of 2500 kiloohms and 1000 kiloohms were connected in series. The voltage across each resistor is calculated using Ohm's law, which is given by equation 1 below.

$$V = I \cdot R \tag{1}$$

In this equation, V represents the voltage across the resistor, I is the current flowing through the resistor, and R is the resistance of the resistor. The resulting voltages across the resistors are as follows:

- The voltage across the 2.5 kiloohm resistor is approximately 17.14 volts.

- The voltage across the 1 kiloohm resistor is approximately 6.86 volts. This voltage was then supplied to the control wire of servo motor 1 of the mini-robot.

Connecting a second servo motor of the mini-robot means repeating the process of reducing the voltage through resistors for the second PWM output of the PW 022 module. The connection between the PW 022 module and the motor of the mini-robot is demonstrated in Figure 20.

Figure 20: Connection between PW 022 and mini-robot [28]

The program was written in Lasal Class 2 and was applied to all three mentioned versions of Salamander 4 to measure the reaction time of the robot to the specific commands. Prior to examining the software program, the next two subsections briefly explain the setup of each version of the experiment.

7.0.1 Setup of Hardware Salamander 4

In this version, Salamander 4 runs on the CP 841 [8] CPU unit, specifically designed for the Salamander 4 operating system. The PW 022 module is directly mounted on the CPU via the S-DIAS bus and communicates over the hard real-time capable Ethernet VARAN with 100 Mbit/s [29]. The setup is visible on Figure 21.

Figure 21: Hardware tree

7.0.2 Setup of QEMU Salamander 4

In the virtualized setup of the experiment, the VCPU functionality of QEMU is used to connect QEMU with the PWM module. In order to achieve that, the PCV 522 VARAN Manager PCI Insert Card [30], which serves as a bridge between the PC and the rest of the setup, needed to be plugged into the PC. The command `lspci -nn` lists all PCI devices along with their vendor and device IDs. In this case, the command for binding the PCV 522 module to the VFIO-PCI driver was `sudo sh -c 'echo "5112 2200" > /sys/bus/pci/drivers/vfio-pci/new_id'` with vendor ID 5112 and device ID 2200. To verify that the device has been successfully bound to the VFIO-PCI driver, `lspci -v` can be used. On top of the binding process, the QEMU script from the previous sections also needed to be modified to include the PCV 522 VARAN Manager PCI Insert Card. This is done by adding `-device vfio-pci,host=03:00.0` to the QEMU script, where `03:00.0` refers to the PCI address of the device, with `03` being the bus number, `00` the device number, and `0` the function number. The final script can be seen in Code 11.

```
1   #!/bin/sh
2
3   if  [ ! -d drive-c/ ]; then
4           echo "Filling drive-c/"
5           mkdir drive-c/
6           tar -C drive-c/ -xf stek-drive-c-image-sigmatek-core2.tar.gz
7   fi
8
9   exec taskset -c 4-5 chrt -f 99 qemu-system-x86_64 -M pc,accel=kvm -kernel
        ./bzImage \
10  -m 2048 -drive
        file=salamander-image-sigmatek-core2.ext4,format=raw,media=disk \
11  -append "console=ttyS0 console=tty1 root=/dev/sda rw panic=1
        sigmatek_lrt.QEMU=1 ip=dhcp rootfstype=ext4 schedstats=enable nohlt
        idle=poll quiet xeno_hal.smi=1 xenomai.smi=1 threadirqs" \
12  -net nic,model=e1000,netdev=e1000 -netdev bridge,id=e1000,br=nm-bridge \
13  -fsdev local,security_model=none,id=fsdev0,path=drive-c -device
        virtio-9p-pci,id=fs0,fsdev=fsdev0,mount_tag=/mnt/drive-C \
14  -device vhost-vsock-pci,guest-cid=3,id=vsock0 \
15  -drive if=pflash,format=qcow2,file=ovmf.code.qcow2 \
16  -object memory-backend-ram,id=ram0,size=4G,prealloc=on \
17  -mem-prealloc -mem-path /dev/hugepages \
18  -device vfio-pci,host=03:00.0 \
19  -no-reboot -nographic
```

Code 11: Include PCI in QEMU script for Salamander 4 virtualization

An additional Varan Connection module VI 021 [31] was required to enable the connection between the PCV 522 module and the PW 022 module that generates the signal to move the mini-robot. This setup is depicted in Figure 22.

Figure 22: Virtualization tree

Essentially, the CP 841 CPU is being emulated by the VCPU functionality of QEMU, allowing the PCV 522 VARAN Manager PCI Insert Card and the Varan Connection module VI 021 to interact as if they were communicating with a physical CPU unit.

7.1 Robotic Application

8 Results

[5] [10] [11] [7] [8] [9] [6] [12] [13] [14] [15] [16] [17] [18] [19] [20] [21] [22] [23] [24] [25] [26] [27] [28] [29] [30] [31] [32] [33] [34] [35] [36] [37] [38] [39] [40] [41] [42] [43] [44] [45] [46] [47] [48] [49] [50] [51] [52]

9 Discussion

10 Summary and Outlook

Bibliography

[1]   Seçkin Canbaz and Gökhan Erdemir. "Performance Analysis of Real-Time and General-Purpose Operating Systems for Path Planning of the Multi-Robot Systems". In: *International Journal of Electrical and Computer Engineering (IJECE)* 12.1 (Feb. 2022), p. 285. ISSN: 2722-2578, 2088-8708. DOI: 10 . 11591 / ijece . v12i1 . pp285 - 292. (Visited on 08/11/2024).

[2]   Jiacun Wang. *Real-Time Embedded Systems*. Hoboken, NJ, USA: Wiley, 2017. ISBN: 978-1-119-42070-5.

[3]   Giorgio C. Buttazzo. *Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and Applications*. Fourth edition. Cham, Switzerland: Springer, 2024. ISBN: 978-3-031-45409-7. DOI: 10.1007/978-3-031-45410-3.

[4]   Giuseppe Lipari, Luigi Palopoli, and Università di Trento. "Real-Time Scheduling: From Hard to Soft Real-Time Systems". In: ().

[5]   pixelart. *SIGMATEK - Komplette Automatisierungssysteme*. https://www.sigmatek-automation.com/de/. (Visited on 03/27/2024).

[6]   *Welcome to the Yocto Project Documentation — The Yocto Project ® 4.3.999 Documentation*. https://docs.yoctoproject.org/. (Visited on 03/27/2024).

[7]   *Xenomai :: Xenomai*. https://xenomai.org/. (Visited on 03/21/2024).

[8]   *CPU-Einheiten - SIGMATEK*. https://www.sigmatek-automation.com/de/produkte/steuerungssysteme/cpu-einheiten/cp-841/. (Visited on 05/27/2024).

[9]   *Engineering Tool LASAL - SIGMATEK*. https://www.sigmatek-automation.com/de/produkte/engineering-tool-lasal/lasal-class/. (Visited on 05/27/2024).

[10] *Trace-Cmd*. https://trace-cmd.org/. (Visited on 03/25/2024).

[11] *KernelShark*. https://kernelshark.org/. (Visited on 03/25/2024).

[12] *QEMU*. https://www.qemu.org/. (Visited on 03/27/2024).

[13] *Realtime:Preempt_rt_versions [Wiki]*. https://wiki.linuxfoundation.org/realtime/preempt_rt_versions. (Visited on 08/05/2024).

[14] *Realtime Kernel Patchset - ArchWiki*. https://wiki.archlinux.org/title/Realtime_kernel_patchset. (Visited on 08/05/2024).

[15] Steven Rostedt and Darren V Hart. "Internals of the RT Patch". In: ().

[16] *What Is Real-Time Linux? Part I*. https://ubuntu.com/blog/what-is-real-time-linux-i. (Visited on 08/05/2024).

[17] Ruhui Ma et al. "Performance Tuning Towards a KVM-based Embedded Real-Time Virtualization System". In: ().

[18] *Linux Process Priorities Demystified*. https://sigma-star.at/blog/2022/02/linux-proc-prios/. (Visited on 06/11/2024).

[19] Kernel Recipes. *Kernel Recipes 2016 - Understanding a Real-Time System (More than Just a Kernel) - Steven Rostedt*. Oct. 2016. (Visited on 06/11/2024).

[20] The Linux Foundation. *Finding Sources of Latency on Your Linux System - Steven Rostedt, VMware*. Sept. 2020. (Visited on 06/11/2024).

[21] The Linux Foundation. *A Checklist for Writing Linux Real-Time Applications - John Ogness, Linutronix GmbH*. Nov. 2020. (Visited on 06/11/2024).

[22] *HOWTO: Build an RT-application - RTwiki*. https://archive.kernel.org/oldwiki/rt.wiki.kernel.org/index.php/HOWTO:_Build_an_RT-application.html. (Visited on 06/11/2024).

[23] *Real-Time Programming with Linux, Part 2: Configuring Linux for Real-Time - Shuhao's Blog*. https://shuhaowu.com/blog/2022/02-linux-rt-appdev-part2.html#f4. (Visited on 06/11/2024).

[24] *KVM/Qemu Virtualization Tuning Guide on Intel® Xeon® Based Systems*. https://www.intel.com/content/www/us/en/developer/articles/guide/kvm-tuning-guide-on-xeon-based-systems.html. (Visited on 06/11/2024).

[25] "Real-Time Performance Tuning Best Practice Guidelines for KVM-Based Virtual Machines". In: (2022).

[26] *MG996R Servo Motor*. https://components101.com/motors/mg996r-servo-motor-datasheet. (Visited on 07/30/2024).

[27] *Digital Output - SIGMATEK*. https://www.sigmatek-automation.com/en/products/io-systems/s-dias/digital-output/pw-022/. (Visited on 07/30/2024).

[28] *MG996R Digital Servo Motor mit Metall Getriebe*. https://www.roboter-bausatz.de/p/mg996r-digital-servo-motor-mit-metall-getriebe. (Visited on 07/31/2024).

[29] *S-DIAS - SIGMATEK*. https://www.sigmatek-automation.com/en/products/io-systems/s-dias/. (Visited on 07/30/2024).

[30] *Controls & HMIs - SIGMATEK*. https://www.sigmatek-automation.com/en/products/accessories/controls-hmis/pcv-522/. (Visited on 07/30/2024).

[31] *Interfaces & Splitters - SIGMATEK*. https://www.sigmatek-automation.com/en/products/real-time-ethernet-varan/interfaces-splitters/vi-021/. (Visited on 07/30/2024).

[32] George K. Adam, Nikos Petrellis, and Lambros T. Doulos. "Performance Assessment of Linux Kernels with PREEMPT_RT on ARM-Based Embedded Devices". In: *Electronics* 10.11 (June 2021), p. 1331. ISSN: 2079-9292. DOI: 10.3390/electronics10111331. (Visited on 05/08/2024).

[33] George K. Adam. "Real-Time Performance and Response Latency Measurements of Linux Kernels on Single-Board Computers". In: *Computers* 10.5 (May 2021), p. 64. ISSN: 2073-431X. DOI: 10.3390/computers10050064. (Visited on 05/16/2024).

[34] S. Brosky and S. Rotolo. "Shielded Processors: Guaranteeing Sub-Millisecond Response in Standard Linux". In: *Proceedings International Parallel and Distributed Processing Symposium*. Nice, France: IEEE Comput. Soc, 2003, p. 9. ISBN: 978-0-7695-1926-5. DOI: 10.1109/IPDPS.2003.1213237. (Visited on 04/18/2024).

[35] Marcello Cinque et al. "Virtualizing Mixed-Criticality Systems: A Survey on Industrial Trends and Issues". In: *Future Generation Computer Systems* 129 (Apr. 2022), pp. 315–330. ISSN: 0167739X. DOI: 10.1016/j.future.2021.12.002. arXiv: 2112.06875 [cs]. (Visited on 03/25/2024).

[36] Daniel Bristot de Oliveira et al. "Demystifying the Real-Time Linux Scheduling Latency". In: ().

[37] Marisol García-Valls, Tommaso Cucinotta, and Chenyang Lu. "Challenges in Real-Time Virtualization and Predictable Cloud Computing". In: *Journal of Systems Architecture* 60.9 (Oct. 2014), pp. 726–740. ISSN: 13837621. DOI: 10.1016/j.sysarc.2014.07.004. (Visited on 03/25/2024).

[38] Zonghua Gu and Qingling Zhao. "A State-of-the-Art Survey on Real-Time Issues in Embedded Systems Virtualization". In: *Journal of Software Engineering and Applications* 05.04 (2012), pp. 277–290. ISSN: 1945-3116, 1945-3124. DOI: 10.4236/jsea.2012.54033. (Visited on 03/25/2024).

[39] "Hard Real Time Linux* Using Xenomai* on Intel® Multi-Core Processors". In: ().

[40] Diogenes Javier Perez et al. "How Real (Time) Are Virtual PLCs?" In: *2022 IEEE 27th International Conference on Emerging Technologies and Factory Automation (ETFA)*. Stuttgart, Germany: IEEE, Sept. 2022, pp. 1–8. ISBN: 978-1-66549-996-5. DOI: 10.1109/ETFA52439.2022.9921545. (Visited on 03/25/2024).

[41] Veronika Kirova et al. "Impact of Modern Virtualization Methods on Timing Precision and Performance of High-Speed Applications". In: *Future Internet* 11.8 (Aug. 2019), p. 179. ISSN: 1999-5903. DOI: 10.3390/fi11080179. (Visited on 03/25/2024).

[42] Jan Kiszka. "Towards Linux as a Real-Time Hypervisor". In: ().

[43] CC Huang, Chan-Hsiang Lin, and Che-Kang Wu. "Performance Evaluation of Xenomai 3". In: *Proceedings of the 17th Real-Time Linux Workshop (RTLWS)*. 2015, pp. 21–22.

[44] Petro Lutsyk, Jonas Oberhauser, and Wolfgang J. Paul. *A Pipelined Multi-Core Machine with Operating System Support: Hardware Implementation and Correctness Proof*. Lecture Notes in Computer Science Theoretical Computer Science and General Issues 9999. Cham: Springer, 2020. ISBN: 978-3-030-43242-3.

[45] HayfaaSubhi Malallah et al. "A Comprehensive Study of Kernel (Issues and Concepts) in Different Operating Systems". In: *Asian Journal of Research in Computer Science* (May 2021), pp. 16–31. ISSN: 2581-8260. DOI: 10.9734/ajrcos/2021/v8i330201. (Visited on 03/25/2024).

[46] Alejandro Masrur et al. "VM-Based Real-Time Services for Automotive Control Applications". In: *2010 IEEE 16th International Conference on Embedded and Real-Time Computing Systems and Applications*. Macau, China: IEEE, Aug. 2010, pp. 218–223. ISBN: 978-1-4244-8480-5. DOI: 10.1109/RTCSA.2010.38. (Visited on 03/25/2024).

[47] Paul E McKenney. "'Real Time' vs. 'Real Fast': How to Choose?" In: ().

[48] Éric Piel et al. "Asymmetric Scheduling and Load Balancing for Real-Time on Linux SMP". In: *Parallel Processing and Applied Mathematics*. Ed. by David Hutchison et al. Vol. 3911. Berlin, Heidelberg: Springer Berlin Heidelberg, 2006, pp. 896–903. ISBN: 978-3-540-34141-3 978-3-540-34142-0. DOI: 10.1007/11752578_108. (Visited on 05/06/2024).

[49] Rui Queiroz, Tiago Cruz, and Paulo Simões. "Testing the Limits of General-Purpose Hypervisors for Real-Time Control Systems". In: *Microprocessors and Microsystems* 99 (June 2023), p. 104848. ISSN: 01419331. DOI: 10.1016/j.micpro.2023.104848. (Visited on 03/25/2024).

[50] Federico Reghenzani, Giuseppe Massari, and William Fornaciari. "The Real-Time Linux Kernel: A Survey on PREEMPT_RT". In: *ACM Computing Surveys* 52.1 (Jan. 2020), pp. 1–36. ISSN: 0360-0300, 1557-7341. DOI: 10.1145/3297714. (Visited on 03/25/2024).

[51] Hobin Yoon, Jungmoo Song, and Jamee Lee. "Real-Time Performance Analysis in Linux-Based Robotic Systems". In: ().

[52] *CPU Units - SIGMATEK*. https://www.sigmatek-automation.com/en/products/control-systems/cpu-units/cp-841/. (Visited on 07/30/2024).

List of Figures

List of Tables

List of Code

List of Abbreviations

CPU   Central Processing Unit

QEMU Quick Emulator

IRQ   Interrupt Request