MASTER THESIS

Thesis submitted in partial fulfillment of the requirements for the degree of Master of Science in Engineering at the University of Applied Sciences Technikum Wien - Degree Program Mechatronics/Robotics

Virtualisierung eines Echtzeit-Betriebssystems zur Steuerung eines Roboters mit Schwerpunkt auf die Einhaltung der Echtzeit

By: Halil Pamuk, BSc

Student Number: 51842568

Supervisor: Sebastian Rauh, MSc. BEng

Wien, June 5, 2024

Declaration

Wien, June 5, 2024   Signature

Kurzfassung

Erstellung einer Echtzeit-Robotersteuerungsplattform unter Verwendung von Salamander OS, Xenomai, QEMU und PCV-521 in der Yocto-Umgebung. Die Plattform basiert auf Salamander OS und nutzt Xenomai für Echtzeit- Funktionen. Dazu muss im ersten Schritt die Virtualisierungsplattform evaluiert werden. (QEMU, Hyper-V, Virtual Box, etc.) Als weiterer Schritt folgt die Anbindung eines Roboters über eine VARAN-Bus Schnittstelle. Das gesamte System wird in der Yocto-Umgebung erstellt und konfiguriert. Das Hauptziel der Arbeit ist es, herauszufinden, wie die Integration von Echtzeit-Funktionen und effizienten Kommunikationssystemen in eine Robotersteuerungsplattform die Reaktionszeit und Zuverlässigkeit von Roboteranwendungen verbessern kann

Schlagworte: Schlagwort1, Schlagwort2, Schlagwort3, Schlagwort4

Abstract

Sections 4.1 and 4.2 demonstrate the inital real-time latency values gathered for bare metal and virtualization.

Contents

1 Introduction

In today's industrial production and automation, robot systems are well established and of crucial importance. Robots must react to their environment and perform time-critical tasks within strict time constraints. Delays or errors can have catastrophic consequences in some cases. Traditional operating systems, such as Windows or Linux, are often not suitable for these types of real-time requirements as they cannot guarantee deterministic execution times. Therefore, real-time operating systems are required that are specifically designed to react to events within fixed time limits and prioritise the execution of high-priority processes.

The core component of an RTOS that enables real-time capabilities is the kernel. The kernel is responsible for managing system resources, scheduling tasks, and ensuring deterministic behavior. It employs preemptive scheduling mechanisms to allow high-priority tasks to preempt lower-priority tasks, ensuring that time-critical tasks are not delayed. The kernel also implements priority-based scheduling algorithms, such as Rate Monotonic Scheduling (RMS) or Earliest Deadline First (EDF), to schedule tasks based on their priorities and timing constraints. Additionally, RTOS kernels are designed to minimize interrupt latency, which is crucial for real-time applications that require immediate response to external events.

In these RTOS systems, task scheduling is based on so-called priority-based preemptive scheduling. Each task in a software application is assigned a priority. A higher priority means that a faster response is required. Preemptive task scheduling ensures a very fast response. Preemptive means that the scheduler can stop a currently running task at any point if it recognizes that another task needs to be executed immediately. The basic rule on which priority-based preemptive scheduling is based is that the task with the highest priority that is ready to run is always the task that must be executed. So if both a task with a lower priority and a task with a higher priority are ready to run, the scheduler ensures that the task with the higher priority runs first. The lower priority task is only executed once the higher priority task has been processed. Real-time systems are usually categorized as either soft or hard real-time systems. The difference lies exclusively in the consequences of a violation of the time limits.

Hard real-time is when the system stops operating if a deadline is missed, which can have catastrophic consequences. Soft real-time exists when a system continues to function even if it cannot perform the tasks within a specified time. If the system has missed the deadline, this has no critical consequences. The system continues to run, although it does so with undesirably lower output quality.

1.1 Application Context

This master's thesis was written at SIGMATEK GmbH & Co KG [1]. SIGMATEK uses its own customized Linux distribution, namely Salamander 4, to be run on their self-manufactured CPUs. Salamander 4 system employs hard real-time with Xenomai 3 and requires a worst latency value between 20 and 50 μs. The goal is to virtualize Salamander 4 and approach the performance of bare metal. Salamander 4 is built with Yocto and virtualized through QEMU/KVM. The details of this operating system are explained in chapter 3.

1.2 State of the art

1.3 Problem and task definition

1.4 Objective

The main objective of this work is to create a real-time robot control platform that integrates Salamander OS, Xenomai, QEMU and PCV-521 in the Yocto environment.

2 Methodology

This section describes in detail all the theoretical concepts and boundary conditions as well as practical methods that contributed to achieving the objectives of this master's thesis.

Trace-cmd was used for tracing the Linux kernel. It can record various kernel events such as interrupts, scheduler decisions, file system activity, function calls in real time. Trace-cmd helped in getting detailed insights into system behaviour and identify reasons for latency [2].

The data that was recorded by trace-cmd was then fed into Kernelshark, which is a graphical front-end tool [3]. It visualizes the recorded kernel trace data in a readable way on an interactive timeline, which facilitated the process of identifying patterns and correlations between events. By further filtering the displayed events according to specific criteria such as processes, event types or time ranges, the latency issues were analyzed.

Real-time operating system capabilities were provided by Xenomai, which is real-time development framework that extends the Linux kernel. It enables low-latency and deterministic execution of time-critical tasks. Xenomai 3 introduces a dual-kernel approach with a real-time kernel coexisting alongside Linux. A key utility within the Xenomai suite is the latency tool, which benchmarks the timer latency - the time it takes for the kernel to respond to timer interrupts or task activations. The tool creates real-time tasks or interrupt handlers and measures the latency between expected and actual execution times [4].

The system configuration is shown in Table 1

Table 1: System configuration

CPU	13th Gen Intel(R) Core(TM) i7-13800H
Memory	2 × 16GB SO-DIMM DDR5-5600 MT/s, 32GB
GPU	NVIDIA RTX A500 Laptop GPU
BIOS	Dell Version 1.12.0
OS	Ubuntu 22.04.4 LTS

Figure 1 is the output of the `lstopo` command and visualizes the hardware nodes of the system, including CPU cores, caches, memory, and I/O devices.


Figure 1: Hardware topology

3 Salamander 4

This chapter briefly describes the Salamander 4 operating system by SIGMATEK.

3.1 Structure

Salamander 4 is the proprietary operating system of SIGMATEK. It is based on Linux version 5.15.94 and integrates Xenomai 3.2, a real-time development environment [4]. Salamander 4 is a 64-bit system, which refers to the x86_64 architecture. The real-time behaviour is achieved through the use of Symmetric Multi-Processing (SMP) and Preemptive Scheduling (PREEMPT). In addition, it uses IRQPIPE to process interrupts in a way that meets the real-time requirements of the system. The output of the command `uname -a` can be observed in code 1.

```
1 root@sigmatek-core2:~# uname -a
2 Linux sigmatek-core2 5.15.94 #1 SMP PREEMPT IRQPIPE Tue Feb 14 18:18:05 UTC
    2023 x86_64 GNU/Linux
```

Code 1: System information

Salamander 4 is powered by SIGMATEK's CP 841 [5] and is comprised of the following software modules:

- **Operating system**: The operating system in a LASAL CPU manages the hardware and software resources of the system. It is provided in a completely PC-compatible manner, working with a standard PC BIOS.

- **Loader**: The loader is a part of the operating system that is responsible for loading programs from executables into memory, preparing them for execution and then executing them.

- **Hardware classes**: Hardware classes in LASAL represent the different types of hardware components that can be controlled by the LASAL CPU. They provide a way to organize and manage the hardware components in a modular and reusable manner. The graphical hardware editor in LASAL allows for a true-to-detail simulation of the actual hardware.

- **Application**: Applications are developed using LASAL CLASS 2 [6], a solution for automation tasks that supports object-oriented programming and design in compliance with IEC 61131-3.

The interfaces between the individual modules are indicated by an arrow in Figure 2.

Figure 2: Structure of Salamander 4 CPU

3.2 Memory Management

For the sake of completeness, Figure 3 displays the memory management of Salamander 4. LRT stands for Lasal Runtime and creates an execution environment where applications developed using the LASAL Class 2 can run, providing defined real-time functions, data types, and other constructs tailored for real-time programming.
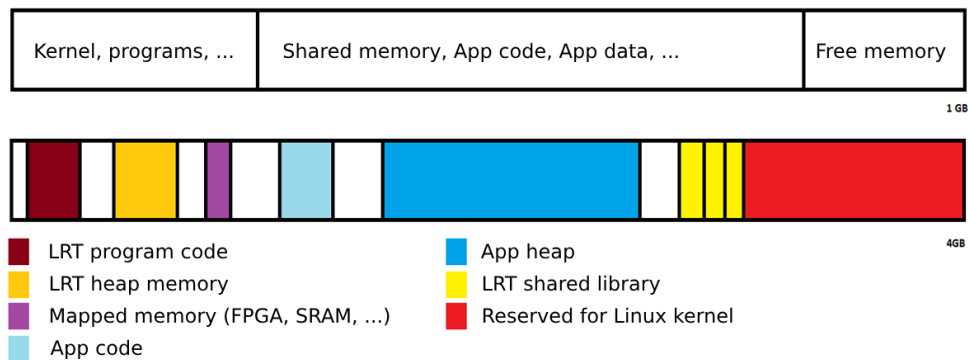
Figure 3: Memory Management

3.3 Xenomai

Xenomai 3 [4] is a real-time framework that offers two paths to real-time performance. The first approach supplements the Linux kernel with a compact real-time core dubbed Cobalt, demonstrated in Figure 4. Cobalt runs side-by-side with Linux, but it handles all time-critical activities like interrupt processing and real-time thread scheduling with higher priority than the regular kernel activities.

Figure 4: Xenomai Cobalt interfaces

The second approach, called Mercury and shown in Figure 5, relies on the real-time capabilities already present in the native Linux kernel. Often, applications require the PREEMPT-RT extension to augment the mainline kernel's real-time responsiveness and minimize jitter, but this isn't mandatory and depends on the application's specific requirements for responsiveness and tolerance for occasional deadline misses.

Figure 5: Xenomai Mercury interfaces

Salamander 4 uses the Cobalt real-time core with the Dovetail extension, which allows the kernel to handle real-time tasks with low latency.

4 Initial Real-Time Latency

As a starting point, initial latency values of both the bare metal and virtualization versions were measured with the latency tool of the xenomai tool suite. Salamander 4 bare metal refers to the proprietary hardware of SIGMATEK used to employ the custom operating system. Salamander 4 virtualization refers to a virtual version of the Salamander 4 hardware platform, achieved through QEMU/KVM. Sections 4.1 and 4.2 specify the details of the measurements for both versions. In the further course, the aim was to bring the latency values of the virtualization closer to those of the hardware and guarantee deterministic and reliable behavior.

## 4.1 Salamander 4 Bare Metal

The output of the command `uname -a` for Salamander 4 bare metal is shown in code 2.

```
1 root@sigmatek-core2:~# uname -a
2 Linux sigmatek-core2 5.15.94 #1 SMP PREEMPT IRQPIPE Tue Feb 14 18:18:05 UTC
    2023 x86_64 GNU/Linux
```

Code 2: Salamander 4 bare metal system information

As a reference point, the latency program was executed on Salamander 4 bare metal for a duration of 10 minutes. The complete command used was `latency -h -g gnuplot.txt -T 600`, which runs the latency measurement tool for 600 seconds and prints histograms of min, avg, max latencies in a Gnuplot-compatible format to the file `gnuplot.txt`. Figure 6 shows the gathered latency values in microseconds.

Figure 6: Latency of Salamander 4 bare metal

Figure 7 depicts the variation in latency over the course of said time. Since the data varies strongly, a logarithmic scale was used for both axes.


Figure 7: Variation in latency of Salamander 4 bare metal

The statistics obtained from this measurement are provided in Table 2. It gives an overview of the average, maximum, minimum latency, and the standard deviation of the latency values in microseconds.

Table 2: Latency statistics of Salamander 4 bare metal in microseconds

Statistic	Value (μs)
Average Latency	4.06
Maximum Latency	10.71
Minimum Latency	2.82
Standard Deviation	0.85

4.2 Salamander 4 Virtualization

In addition to providing Salamander 4 on its own hardware, SIGMATEK has also developed a virtualised version of this operating system. It was developed using Yocto, an open-source project that allows customised Linux distributions to be created for embedded systems [7]. Upon generating the necessary files, Yocto provides a QEMU folder with the following components shown in code 3. QEMU is the environment in which the virtualization runs, as it is an open-source tool for hardware virtualization [8].

```
1    sigma_ibo@localhost:~/Desktop/salamander-image$ ls -1
2    bzImage
3    drive-c
4    ovmf.code.qcow2
5    qemu_def.sh
6    salamander-image-sigmatek-core2.ext4
7    stek-drive-c-image-sigmatek-core2.tar.gz
8    vmlinux
```

Code 3: Contents of QEMU folder for Salamander 4

With the help of the script depicted in code 4, Salamander 4 is started together with the necessary hardware components in the QEMU environment. This makes it possible to run Salamander 4 on a variety of host systems, regardless of the specific hardware of the host. The following is a description of the components used for the virtualization of Salamander 4:

- **bzImage**: Compressed Linux kernel image that is loaded by QEMU at system start. "bz" stands for big-zipped.

- **drive-c**: Directory serving as C drive for QEMU system, created and filled by qemu_def.sh script.

- **ovmf.code.qcow2**: Firmware file for QEMU that enables UEFI boot process. `OVMF` stands for Open Virtual Machine Firmware, `qcow2` is a format for disk image files used by QEMU, it stands for "QEMU Copy On Write version 2".

- **qemu_def.sh**: Shell script that starts QEMU with required parameters to boot Salamamder 4 OS. It is described in Code 4.

- **salamander-image-sigmatek-core2.ext4**: Disk image of the Salamander 4 OS for the Sigmatek Core 2 platform. It uses the `ext4` file system and serves as the root file system in the QEMU virtual machine, acting as the virtual hard drive.

- **stek-drive-c-image-sigmatek-core2.tar.gz**: Compressed tarball containing a pre-configured environment for the Salamander 4 OS. It is unpacked and sets up the `drive-c/` directory with system and log files in the `qemu_def.sh` script.

- **vmlinux**: Uncompressed Linux kernel image, typically used for debugging.

The inital QEMU script after the custom Yocto build and the starting point for this work is shown in Code 4. This script is used to start QEMU with required parameters to boot Salamamder 4 OS. It will be adjusted in chapter 5 in order to accompany real-time performance tunings.

```
1   #!/bin/sh
2
3   if  [ ! -d drive-c/ ]; then
4         echo "Filling drive-c/"
5         mkdir drive-c/
6         tar -C drive-c/ -xf stek-drive-c-image-sigmatek-core2.tar.gz
7   fi
8
9   exec qemu-system-x86_64 -M pc,accel=kvm -kernel ./bzImage \
10  -m 2048 -drive
        file=salamander-image-sigmatek-core2.ext4,format=raw,media=disk \
11  -append "console=ttyS0 console=tty1 root=/dev/sda rw panic=1
        sigmatek_lrt.QEMU=1 ip=dhcp rootfstype=ext4" \
12  -net nic,model=e1000,netdev=e1000 -netdev bridge,id=e1000,br=nm-bridge \
13  -fsdev local,security_model=none,id=fsdev0,path=drive-c -device
        virtio-9p-pci,id=fs0,fsdev=fsdev0,mount_tag=/mnt/drive-C \
14  -device vhost-vsock-pci,guest-cid=3,id=vsock0 \
15  -drive if=pflash,format=qcow2,file=ovmf.code.qcow2 \
16  -no-reboot -nographic
```

Code 4: QEMU script for starting Salamander 4 virtualization

This script is run on a generic Ubuntu 22.04.4 system, as mentioned previously in chapter 2. The kernel version and other details are presented in Code 5, using the `uname -a` command.

```
1 root@sigmatek-core2:~# uname -a
2 Linux sigma-ibo 6.5.0-35-generic #35~22.04.1-Ubuntu SMP PREEMPT_DYNAMIC Tue
     May  7 09:00:52 UTC 2 x86_64 x86_64 x86_64 GNU/Linux
```

Code 5: Ubuntu 22.04.4 system information

Measuring the latency of the Salamander 4 virtualization with the default QEMU script in Code 4 and no further adjustments for 10 minutes, the following latency values in Figure 8 were collected.

Figure 8: Latency with default settings

Figure 9 depicts the variation in latency over the course of said time with default settings.

Figure 9: Variation in latency with default settings

The statistics obtained from this measurement are provided in Table 3. It gives an overview of the average, maximum, minimum latency, and the standard deviation of the latency values in microseconds.

Table 3: Latency statistics of default Salamander 4 virtualization in microseconds

Statistic	Value (µs)
Average Latency	46.22
Maximum Latency	707.62
Minimum Latency	15.59
Standard Deviation	52.13

Comparing these values to those of bare metal, it is evident that there is a significant initial gap in the statistics. A maximum latency of 707.62 µs is not tolerable for the system and needs to be tuned.

5 Real-Time Performance Tuning

In this chapter, the significant initial gap in latency statistics between the virtualized system and the bare metal system is tackled. For this reason, an extensive tuning process is carried out. This involves configurations spanning the BIOS, kernel, host OS, QEMU-KVM virtualization layer, and the Salamander 4 OS itself. The individual configurations will be discussed in detail and the modifications will be justified with a clear explanation. The goal is to bring the latency of the virtualized system closer to that of the bare metal system, thereby ensuring deterministic behavior under real-time constraints.

5.1 BIOS Configurations

BIOS stands for Basic Input/Output System. It abstracts the hardware and enables basic functions of a computer during the booting process, such as starting the operating system and loading other software. Since the BIOS is embedded very deep, its configuration can significantly influence the real-time performance of the system. Table 4 illustrates the specific BIOS settings that have been adjusted for the purpose of real-time performance.

Table 4: BIOS Configurations

Option	Status
Hyper Threading	Disabled
Intel SpeedStep®	Disabled
Intel® Speed Shift Technology	Disabled
C States	Disabled
VT-d	Enabled

In the following, these settings along with their impact on system latency are briefly described.

- **Hyper Threading**: Hyper-Threading allows CPUs to process two threads simultaneously instead of just one. When it is enabled on the host, this allows the parallelisation of tasks and increases performance. However, in a real-time system like the guest Salamander 4, this can lead to increased latencies due to contention between threads. In order to ensure more deterministic behavior in the guest, it is disabled on the host.

- **Intel SpeedStep**: This technology dynamically adjusts the clock speed of the CPU based on workload. These dynamic changes can lead to unpredictable latencies in a real-time system. It is also disabled on the host to maintain a constant CPU speed.

- **Intel® Speed Shift Technology**: Similar to SpeedStep, Speed Shift allows the processor to directly control its frequency and voltage. This can lead to unpredictable latencies, too. Hence, it is also disabled on the host.

- **C States**: These are low-power idle states where the clock frequency and voltage of the CPU are reduced. Transitioning between C-states can cause variable latencies. To prevent this from happening, C-states are disabled on the host.

- **VT-d**: Direct access to physical devices from within virtual machines is possible when VT-d is enabled on the host. This can help reduce latencies associated with I/O operations in the virtual machine. It is therefore enabled on the host.

5.2 Kernel Configurations

The kernel command-line parameters are shown in Code 6 below. Some of them are essential for real-time performance.

```
1        GRUB_CMDLINE_LINUX="isolcpus=4 rcu_nocbs=4 nohz_full=4
         default_hugepagesz=1G hugepagesz=1G hugepages=8 intel_iommu=on
         rdt=l3cat nmi_watchdog=0 idle=poll clocksource=tsc tsc=reliable
         noht audit=0 skew_tick=1 intel_pstate=disable intel.max_cstate=0
         intel_idle.max_cstate=0 processor.max_cstate=0
         processor_idle.max_cstate=0 nosoftlockup nohz=on no_timer_check
         nospectre_v2 spectre_v2_user=off kvm.kvmclock_periodic_sync=N
         kvm_intel.ple_gap=0 irqaffinity=0"
```

Code 6: Kernel Configuration

In the following, these settings along with their impact on system latency are briefly described.

- **isolcpus=4**: Isolates CPU 4 from the general scheduler, meaning no process will be scheduled to run on this CPU unless it is explicitly assigned. CPU isolation is explained in detail in section 5.3.1.

- **rcu_nocbs=4**: The Linux kernel uses a synchronization mechanism called RCU, or read-copy-update. It lets writers update the data in a way that guarantees readers will always see the same version while enabling multiple readers to access shared data without locks. The RCU subsystem uses callback functions that need to be invoked once readers are done with the data they accessed. By default, these callbacks are handled by the CPUs that executed the read-side critical sections. This parameter offloads RCU callback handling from CPU 4 to other CPUs. CPU 4 remains dedicated to high-priority tasks which helps in reducing latency.

- **nohz_full=4**: Makes CPU core 4 "tickless", meaning the kernel tries to avoid sending periodic scheduling-clock interrupts to the CPU when there are no runnable tasks. This lowers latency by reducing unnecessary wake-ups but may increase power consumption because the CPU is not able to enter a low-power state when idle. Additionally, timer interrupts cannot be fully eliminated because certain events, such as incoming interrupts or task activations, can still cause the kernel to send timer interrupts to the tickless CPU.

5.3 Host OS Configurations

5.3.1 CPU isolation

Isolating CPUs involves removing all user-space threads and unbound kernel threads since bound kernel threads are tied to specific CPUs and hence cannot be moved. Also, modifying the `proc/irq/IRQ_NUMBER/smp_affinity` property of each Interrupt IRQ_NUMBER in the system is part of this process, as described later in section 5.3.2.

For CPU isolation, the `isolcpus` function was used to isolate a peroformance CPU from the general scheduling algorithms of the operating system. This means that the isolated CPUs will not be used for regular task scheduling, allowing them to be dedicated for the real-time specific tasks or processes. However, the isolcpus function only isolates at the user level and does not affect kernel tasks. Consequently, these kernel tasks and interrupts can still utilize the CPU [9]. Output 7 shows the user and kernel tasks that run on CPU 4. After the isolation, user tasks other than the QEMU process have been removed from running on this CPU. Only few critical kernel threads that are tied to this CPU still take CPU time.

```
1    sigma_ibo@sigma-ibo:~$ cat /sys/devices/system/cpu/isolated
2    4
3    sigma_ibo@sigma-ibo:~$ ps axHo psr,pid,lwp,args,policy,nice,rtprio |
         awk '$1 == 4'
4    4      38      38 [cpuhp/4]                   TS    0     -
5    4      39      39 [idle_inject/4]             FF    -     50
6    4      40      40 [migration/4]               FF    -     99
7    4      41      41 [ksoftirqd/4]               TS    0     -
8    4      43      43 [kworker/4:0H-events_highpr TS   -20    -
9    4     505     505 [irq/189-iwlwifi:queue_1]   FF    -     50
10   4     519     519 [irq/203-iwlwifi:exception] FF    -     50
11   4    3008    3008 [kworker/4:1H-kblockd]      TS   -20    -
12   4  177779  177779 [kworker/4:2-events]        TS    0     -
13   4  448807  448807 [kworker/4:1-events]        TS    0     -
```

Code 7: User and Kernel Tasks

5.3.2 Interrupt Requests Handling

Once the CPUs were isolated, Interrupt Requests handling was the next step. Interrupt Requests are used to send a signal to the CPU, prompting it to 'interrupt' its current task and divert its attention to another task. This allows hardware devices to communicate with the CPU through frequent context switches, which can lead to performance degradation, especially in high-performance computing or real-time scenarios. To mitigate this, the IRQs needed to be removed from the isolated CPU. This was done by manipulating a file in the `proc` filesystem, namely `/proc/irq/<IRQ>/smp_affinity`. The value in the `smp_affinity` file is a bitmask in hexadecimal format. Each bit in this mask corresponds to a CPU in the system. The least significant bit (LSB) on the right corresponds to the first CPU (CPU0), and the significance increases towards the left until CPU19. In a system with 14 CPUs, if every CPU was reserved for one IRQ, the value for `smp_affinity` would be 3FFF. The script in code 8 was written to check and log the distribution of Interrupt Requests across each CPU in Salamander 4.

```bash
1    #!/bin/bash
2    # Check if a command-line argument is provided
3    if [ -z "$1" ]; then
4        echo "Please provide a CPU number as a command-line argument."
5        exit 1
6    fi
7    # Get the CPU number from the command-line argument
8    CPU=$1
9    # Initialize an empty array to store the IRQ numbers
10   IRQs=()
11   for IRQ in /proc/irq/*; do
12       if [ -f "$IRQ/smp_affinity" ]; then
13           # Read the current smp_affinity
14           AFFINITY=$(cat "$IRQ/smp_affinity")
15           # Check if the bit for the current CPU is set
16           if (( (0x$AFFINITY & (1 << CPU)) != 0 )); then
17               # Add the IRQ number to the array
18               IRQs+=("${IRQ#/proc/irq/}")
19           fi
20       fi
21   done
22   # Sort the array
23   IFS=$'\n' sorted=($(sort -n <<<"${IRQs[*]}"))
24   # Print the CPU number
25   echo "CPU $CPU IRQ affinity:"
26   # Print the sorted IRQ numbers on separate lines
27   for irq in "${sorted[@]}"; do
28       echo "$irq"
29   done
```

Code 8: Check distribution of Interrupt Requests across each CPU

By changing the values of the `smp_affinity` files of the respective IRQs, the assignment of

IRQs was controlled so that they would not be handled by the isolated CPU. This reduced the interruptions caused by IRQs and the isolated CPUs were able to focus more on their assigned tasks.

5.3.3 Disable dynamic frequency scaling

5.3.4 Disable RT throttling

5.3.5 No unexpected RT processes are running on your system

5.3.6 IRQ affinity

5.3.7 RCU CPU offloading

5.3.8 Suppress rcu cpu stall

5.3.9 Maybe?

5.3.10 Start QEMU normally and give all QEMU threads rt-priority

5.3.11 Kill all running user processes

5.3.12 Set CPU Affinity for systemd services

5.3.13 Cache Isolation for CPU and GPU

5.3.14 Set CPU Affinity of IRQ thread to CPU 0

5.3.15 Set Device Driver Work Queue to CPU 0

5.3.16 Disable Machine Check

5.3.17 Stop Certain Services

5.4 QEMU-KVM Configurations

5.4.1 Tune lapic timer advance

5.4.2 Set QEMU options for real-time VM

5.4.3 Set CPU affinity and scheduling policy of QEMU CPU threads

5.4.4 Passthrough PCI devices into the VM

5.5 Guest OS Configurations

6 KVM exit reasons

6.1 APIC_WRITE

6.2 HLT

6.3 EPT_MISCONFIG

6.4 PREEMPTION_TIMER

6.5 EXTERNAL_INTERRUPT

6.6 IO_INSTRUCTION

6.7 EOI_INDUCED

6.8 EPT_VIOLATION

6.9 PAUSE_INSTRUCTION

6.10 CPUID

6.11 MSR_READ

7 Real-Time Robotic Application

7.1 VARAN

7.2 Robotic Application

8 Results

[1] [2] [3] [4] [5] [6] [7] [8] [9] [10] [11] [12] [13] [14] [15] [16] [17] [18] [19] [20] [21] [22] [23] [24] [25] [26] [27] [28] [29] [30]

9 Discussion

10 Summary and Outlook

Bibliography

[1] pixelart. *SIGMATEK - Komplette Automatisierungssysteme*. https://www.sigmatek-automation.com/de/. (Visited on 03/27/2024).

[2] *Trace-Cmd*. https://trace-cmd.org/. (Visited on 03/25/2024).

[3] *KernelShark*. https://kernelshark.org/. (Visited on 03/25/2024).

[4] *Xenomai :: Xenomai*. https://xenomai.org/. (Visited on 03/21/2024).

[5] *CPU-Einheiten - SIGMATEK*. https://www.sigmatek-automation.com/de/produkte/steuerungssysteme/cpu-einheiten/cp-841/. (Visited on 05/27/2024).

[6] *Engineering Tool LASAL - SIGMATEK*. https://www.sigmatek-automation.com/de/produkte/engineering-tool-lasal/lasal-class/. (Visited on 05/27/2024).

[7] *Welcome to the Yocto Project Documentation — The Yocto Project ® 4.3.999 Documentation*. https://docs.yoctoproject.org/. (Visited on 03/27/2024).

[8] *QEMU*. https://www.qemu.org/. (Visited on 03/27/2024).

[9] Ruhui Ma et al. "Performance Tuning Towards a KVM-based Embedded Real-Time Virtualization System". In: ().

[10] Chan-Hsiang Lin and Che-Kang Wu. "Performance Evaluation of Xenomai 3". In: ().

[11] S. Brosky and S. Rotolo. "Shielded Processors: Guaranteeing Sub-Millisecond Response in Standard Linux". In: *Proceedings International Parallel and Distributed Processing Symposium*. Nice, France: IEEE Comput. Soc, 2003, p. 9. ISBN: 978-0-7695-1926-5. DOI: 10.1109/IPDPS.2003.1213237. (Visited on 04/18/2024).

[12] Marcello Cinque et al. "Virtualizing Mixed-Criticality Systems: A Survey on Industrial Trends and Issues". In: *Future Generation Computer Systems* 129 (Apr. 2022), pp. 315–330. ISSN: 0167739X. DOI: 10.1016/j.future.2021.12.002. arXiv: 2112.06875 [cs]. (Visited on 03/25/2024).

[13] Marisol García-Valls, Tommaso Cucinotta, and Chenyang Lu. "Challenges in Real-Time Virtualization and Predictable Cloud Computing". In: *Journal of Systems Architecture* 60.9 (Oct. 2014), pp. 726–740. ISSN: 13837621. DOI: 10.1016/j.sysarc.2014.07.004. (Visited on 03/25/2024).

[14] Zonghua Gu and Qingling Zhao. "A State-of-the-Art Survey on Real-Time Issues in Embedded Systems Virtualization". In: *Journal of Software Engineering and Applications* 05.04 (2012), pp. 277–290. ISSN: 1945-3116, 1945-3124. DOI: 10.4236/jsea.2012.54033. (Visited on 03/25/2024).

[15] "Hard Real Time Linux* Using Xenomai* on Intel® Multi-Core Processors". In: ().

[16] Diogenes Javier Perez et al. "How Real (Time) Are Virtual PLCs?" In: *2022 IEEE 27th International Conference on Emerging Technologies and Factory Automation (ETFA)*. Stuttgart, Germany: IEEE, Sept. 2022, pp. 1–8. ISBN: 978-1-66549-996-5. DOI: 10.1109/ETFA52439.2022.9921545. (Visited on 03/25/2024).

[17] Veronika Kirova et al. "Impact of Modern Virtualization Methods on Timing Precision and Performance of High-Speed Applications". In: *Future Internet* 11.8 (Aug. 2019), p. 179. ISSN: 1999-5903. DOI: 10.3390/fi11080179. (Visited on 03/25/2024).

[18] Jan Kiszka. "Towards Linux as a Real-Time Hypervisor". In: ().

[19] Petro Lutsyk, Jonas Oberhauser, and Wolfgang J. Paul. *A Pipelined Multi-Core Machine with Operating System Support: Hardware Implementation and Correctness Proof*. Lecture Notes in Computer Science Theoretical Computer Science and General Issues 9999. Cham: Springer, 2020. ISBN: 978-3-030-43242-3.

[20] HayfaaSubhi Malallah et al. "A Comprehensive Study of Kernel (Issues and Concepts) in Different Operating Systems". In: *Asian Journal of Research in Computer Science* (May 2021), pp. 16–31. ISSN: 2581-8260. DOI: 10.9734/ajrcos/2021/v8i330201. (Visited on 03/25/2024).

[21] Alejandro Masrur et al. "VM-Based Real-Time Services for Automotive Control Applications". In: *2010 IEEE 16th International Conference on Embedded and Real-Time Computing Systems and Applications*. Macau, China: IEEE, Aug. 2010, pp. 218–223. ISBN: 978-1-4244-8480-5. DOI: 10.1109/RTCSA.2010.38. (Visited on 03/25/2024).

[22] Paul E McKenney. "'Real Time' vs. 'Real Fast': How to Choose?" In: ().

[23] Éric Piel et al. "Asymmetric Scheduling and Load Balancing for Real-Time on Linux SMP". In: *Parallel Processing and Applied Mathematics*. Ed. by David Hutchison et al. Vol. 3911. Berlin, Heidelberg: Springer Berlin Heidelberg, 2006, pp. 896–903. ISBN: 978-3-540-34141-3 978-3-540-34142-0. DOI: 10.1007/11752578_108. (Visited on 05/06/2024).

[24] Rui Queiroz, Tiago Cruz, and Paulo Simões. "Testing the Limits of General-Purpose Hypervisors for Real-Time Control Systems". In: *Microprocessors and Microsystems* 99 (June 2023), p. 104848. ISSN: 01419331. DOI: 10.1016/j.micpro.2023.104848. (Visited on 03/25/2024).

[25] Federico Reghenzani, Giuseppe Massari, and William Fornaciari. "The Real-Time Linux Kernel: A Survey on PREEMPT_RT". In: *ACM Computing Surveys* 52.1 (Jan. 2020), pp. 1–36. ISSN: 0360-0300, 1557-7341. DOI: 10.1145/3297714. (Visited on 03/25/2024).

[26] Hobin Yoon, Jungmoo Song, and Jamee Lee. "Real-Time Performance Analysis in Linux-Based Robotic Systems". In: ().

[27] George K. Adam, Nikos Petrellis, and Lambros T. Doulos. "Performance Assessment of Linux Kernels with PREEMPT_RT on ARM-Based Embedded Devices". In: *Electronics* 10.11 (June 2021), p. 1331. ISSN: 2079-9292. DOI: 10.3390/electronics10111331. (Visited on 05/08/2024).

[28] George K. Adam. "Real-Time Performance and Response Latency Measurements of Linux Kernels on Single-Board Computers". In: *Computers* 10.5 (May 2021), p. 64. ISSN: 2073-431X. DOI: 10.3390/computers10050064. (Visited on 05/16/2024).

[29] Daniel Bristot de Oliveira et al. "Demystifying the Real-Time Linux Scheduling Latency". In: ().

[30] "Real-Time Performance Tuning Best Practice Guidelines for KVM-Based Virtual Machines". In: (2022).

List of Figures

List of Tables

List of Code

List of Abbreviations

CPU   Central Processing Unit

QEMU Quick Emulator

IRQ   Interrupt Request

A Anhang A

B Anhang B